



DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics: Computational Science and
Engineering

Thesis title

Ravil Dorozhinskii





DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics: Computational Science and
Engineering

Thesis title

Titel der Abschlussarbeit

Author:	Ravil Dorozhinskii
Supervisor:	Supervisor
Advisor:	Advisor
Submission Date:	Submission date



I confirm that this master's thesis in informatics: computational science and engineering is my own work and I have documented all sources and material used.

Munich, Submission date

Ravil Dorozhinskii

Acknowledgments

Abstract

Parallel linear equation solvers are one of the most important components determining the scalability and efficiency of many supercomputing applications.

reformulate
the state-
ment

In many HPC applications, the optimized single-node performance is one of the most crucial building blocks for the optimization of probabilistic analysis or hybrid MPIOpenMP implementation.

Nomenclature

MPI Message Passing Interface library

OpenMP Open Multi-Processing library

BLAS Basic Linear Algebra Subprograms

LAPACK Linear Algebra Package

ScaLAPACK Scalable Linear Algebra Package

GEMM General Matrix-matrix Multiplication (BLAS subroutine)

TRSM Triangular Solver with Multiple right-hand sides (LAPACK subroutine)

GETRF General Triangular Factorization (LAPACK subroutine)

NUMA non-uniform memory access

PETSc Portable, Extensible Toolkit for Scientific Computation

MUMPS MULTifrontal Massively Parallel sparse direct Solver

WSMP Watson Sparse Matrix Package

PDE Partial Differential Equation

SPD Symmetric Positive Definiten

PE Processing Element

GRS Die Gesellschaft für Anlagen- und Reaktorsicherheit (GRS) gGmbH

LRZ Leibniz-Rechenzentrum (Leibniz Supercomputing Centre of the Bavarian Academy of Sciences and Humanities)

HW1 Hardware installed on the GRS cluster

HW2 Hardware installed on a LRZ CoolMUC-2 Linux cluster

Contents

Acknowledgments	iii
Abstract	iv
1. Solver selection and configuration	1
1.1. Problem statement	1
1.2. Matrix sets, hardware, compilers and MPI library	2
1.3. Direct dense methods	5
1.4. Iterative methods	9
1.5. Direct sparse methods	15
1.6. Hybrid method	31
1.7. Choice of a Sparse Direct Solver Library	33
1.8. Review of MUMPS Library	38
1.9. Choice of Fill Reducing Reordering	42
1.10. MUMPS: Process Pinning	43
1.11. Choice of BLAS Library	45
1.12. MPI-OpenMP Tuning of MUMPS Library	49
1.13. Parallel Performance of MUMPS in case of Small Systems of Equations	53
1.14. Conclusions	53
1.15. Outlook	53
Appendices	54
A. Choice of a Sparse Direct Solver Library	55
List of Figures	57
List of Tables	59
Bibliography	60

1. Solver selection and configuration

1.1. Problem statement

Numerical integration of time dependent partial differential equations can be achieved using different numerical schemes, namely: Runge-Kutta methods. In numerical analysis, the Runge-Kutta methods are a family of implicit and explicit iterative methods [Wik18c]. Implicit and explicit methods have their advantages and disadvantages which can be found in the corresponding literature.

from
PDEs dis-
cretiza-
tion to
ODE

In general, implicit methods are robust in case of stiff equations, however, on another hand they are more expensive in terms of computational cost since they introduce a non-linear part that has to be solve using the Newton's method as an example. There is also a class of integration methods called semi-implicit that were designed to reduce the cost of implicit ones and be able to cope with stiff equations. But in all cases numerical integration boils down to solution of a couple systems of linear equations of a type:

$$Ax = b \quad (1.1)$$

where $A \in \mathbb{R}^{N \times N}$ is an invertible square matrix, $b \in \mathbb{R}^N$ represents the right-hand side, and $x \in \mathbb{R}^N$ is the solution vector. At this point we refer to a system of linear equations as just a system and we will use these two terms interchangeably.

In some cases, especially in case of implicit or semi-implicit methods, we have to compute several systems to perform one step of numerical integration. Thus, solution of a system 1.1 is the computational core of any time integration scheme. In fact, it is the most expensive part and it is primary source of code optimization.

There are three major ways to solve a system of linear equations, namely: direct dense methods, sparse direct methods and iterative methods.

In the following sections 1.3, 1.4, 1.5 we are briefly going to review all three types of numerical solvers as well as their advantages and disadvantages and some aspects

of their parallel implementations with a strong focus on their strong scaling behavior. Strong scaling is important for this research since we want to find a solver and its configuration to solve a certain system with a fixed size as fast as possible. Additionally to that, due to a large number of numerical integration steps, we want to find a robust solver to avoid any application crash during simulations.

At this point we can set requirements for a solver that we are looking for:

- robustness
- numerical stability
- parallel efficiency
- open source licenses

1.2. Matrix sets, hardware, compilers and MPI library

During the study, we are using two matrix sets: GRS and SuiteSparse. In our case, the SuiteSparse matrix set was, in fact, few matrices downloaded from SuiteSparse Matrix Collection [DH11], [DGL89]. We tried to choose different matrices from the collection with respect to both the number of equations n in a system and ratio R between the number of non-zero elements nnz and the number of equations.

NAME:
experimental
setup

To generate GRS matrix set, we ran the most common GRS simulations in ATHLET and stopped the simulations somewhere in the middle saving corresponding shifted Jacobian matrices in the PETSc binary format.

describe
what
ATHLET
is

Tables 1.1 and 1.2 shows matrix properties of both matrix sets.

describe
what a
shifted
Jacobian
matrix is

From now, we are going to introduce and use a definition of **skinny sparse matrices**. It is matrices with relatively low ratio R i.e. less than 15.

The objective of this study was to find and configure a solver which could fulfill all requirements listed above for the GRS matrix set. From time to time, we used the SuiteSparse set for comparison reasons. It seems to us that GRS matrix set was different

add spar-
sity plots
to ap-
pendix

add spar-
sity plots
to ap-
pendix

1. Solver selection and configuration

Name	n	nnz	nnz / n
pwr-3d	6009	117897	19.6200
cube-5	9325	117897	6.0568
cube-64	100657	1388993	13.7993
cube-645	1000045	13906057	13.9054
k3-2	130101	787997	6.0568
k3-18	1155955	7204723	6.2327

Table 1.1.: GRS matrix set

Name	n	nnz	nnz / n	Field
cant	62451	4007383	64.1684	2D/3D Problem
consph	83334	6010480	72.1251	2D/3D Problem
CurlCurl_3	1219574	13544618	11.1060	Model Reduction Problem
Geo_1438	1437960	63156690	43.9210	2D/3D Problem
memchip	2707524	13343948	4.9285	Circuit Simulation Problem
nlpkkt80	1062400	28192672	26.5368	Optimization Problem
PFlow_742	742793	37138461	49.9984	2D/3D Problem
pkustk10	80676	4308984	53.4110	Structural Problem
torso3	259156	4429042	7.0903	2D/3D Problem
x104	108384	8713602	80.3956	Structural Problem

Table 1.2.: SuiteSparse matrix set

because of some specifics of 1D pipeline discretization. For example, GRS matrices are skinny and blocked, where each block is a small, approximately 3-by-3, dense matrix. Additionally, some rows can contain only one element i.e. on the diagonal. It is a result of dynamic pipeline switching of a reactor cooling system. As we will see in section 1.5, sparse direct solvers can be quite sensitive to the sparse structure of a matrix.

We used different hardware to measure performance of different solvers. The first machine was the GRS cluster (HW1) which was our main target. We used a LRZ CoolMUC-2 Linux cluster (HW2) every time when we got some ambiguous results in order to check whether a problem was hardware or software specific. Table 1.3 shows a single node specification of both machines.

We decided to stick to the OpenMPI library which is an implementation of the MPI standard. OpenMPI is an open-source project and it is very well documented. The

	HW1 (GRS)	HW2 (LRZ Linux)
Architecture	x86_64	x86_64
CPU(s)	20	28
On-line CPU(s) list	0-19	0-27
Thread(s) per core	1	1
Core(s) per socket	10	14
Socket(s)	2	2
NUMA node(s)	2	4
Model	62	63
Model name	E5-2680 v2	E5-2697 v3
Stepping	4	2
CPU MHz	1200.0	2036.707
Virtualization	VT-x	VT-x
L1d cache	32K	32K
L1i cache	32K	32K
L2 cache	256K	256K
L3 cache	25600K	17920K
NUMA node0 CPU(s)	0-9	0-6
NUMA node1 CPU(s)	10-19	7-13
NUMA node2 CPU(s)	-	14-20
NUMA node3 CPU(s)	-	21-27

Table 1.3.: Hardware specification

library has many options for processes pinning which was a crucial part in our study because we had to deal with multi-socket machines which, in turn, had multiple NUMA domains.

To make process pinning deterministic, we developed a python script which automatically generated rank-files based on the number of MPI processes, OpenMP threads per MPI process, the maximum number of processing elements and the number of NUMA domains. The scrip always leaves appropriate gaps between MPI processes to allow each process to fork the corresponding number of threads within a parallel region.

A rank-file specifies explicit mapping between MPI processes (ranks) and actual processing elements (cores) within a machine. The script has two modes, namely: *spread* and *close*. Given a certain number of ranks, the spread mode tries to distribute them as spread as possible across available NUMA domains in round-robin fashion.




Figure 1.1.: An example of process pinning of 5 MPI processes with 2 OpenMP threads per rank in case of HW1

In contrast to the spread strategy, the close mode groups ranks as close as possible to keep the maximum number of ranks within a single NUMA domain. Figure 1.1 shows an example of how these two modes work in case of 5 MPI ranks, 2 OpenMP threads per rank, on a compute node equipped with 20 cores and 2 NUMA domains (HW1).

We chose **Intel 18.2** compiler for this study as the newest and the most efficient Intel compiler at the time of writing.

place for figure 1.1

1.3. Direct dense methods

Direct methods have very good numerical stability properties. They do not require to modify a system of equations as it is necessary for iterative methods. Sometimes we only need to permute some rows and columns in order to avoid small absolute values along the matrix diagonal, which can have their negative effect on numerical accuracy, in case of direct methods. However, this operation can be considered relatively cheap.

MUST BE ONLY 2D Scalapack EXAMPLE HERE

On another hand, the computational complexity of $O(n^3)$ and storage requirements of $O(n^2)$ make direct methods not suitable for computation of large systems with more than $\sim 10^3$ number of equations.

Another good property of direct dense solvers is they have direct memory access and they are based on dense linear algebra subroutines. As a result, implementation of these methods can exploit such programming techniques like cache blocking, tiling, data prefetching, utilization of hardware vector units and so on. All of these together make it possible to achieve more than 75% hardware performance of modern CPUs [GG06] due to a high ratio of floating point operations per memory access.

A general way to solve a system of a type 1.1 is to perform LU factorization where a matrix A can be viewed as a product of a lower and upper triangular matrices.

$$Ax = LUx = b \quad (1.2)$$

Having computed matrices L and U , two additional steps, forward and backward substitutions, are required to compute the solution vector x .

$$Ly = b \quad (1.3)$$

$$Ux = y \quad (1.4)$$

The main idea of parallel LU decomposition is based on block structure of a matrix A . Any given matrix A can be represented as a combination of sub-matrices A_{ij} .

$$A = \left(\begin{array}{c|c|c} A_{11} & A_{12} & A_{13} \\ \hline A_{21} & A_{22} & A_{23} \\ \hline A_{31} & A_{32} & A_{33} \end{array} \right)$$

This allows to develop algorithms that can compute LU decomposition of the full matrix A in parallel using sub-matrices with three main routines i.e. matrix-matrix multiply, triangular solve with multiple right hand sides and the unblock LU factorization for operation within a block column [AD90].

There exist three well known parallel approaches based on left, right and Crout decomposition forms. All of these three approaches have similar overall performance, with a slight advantage to the right-looking and Crout variants [AD90].

Let's consider right-looking algorithm as an example. The algorithm can be written as following:

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} = \begin{pmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{pmatrix} \cdot \begin{pmatrix} U_{11} & U_{12} \\ 0 & U_{22} \end{pmatrix} \quad (1.5)$$

The algorithm starts with an assumption that both matrices L_{11} and U_{11} have already been computed.

Given equation 1.5, we can write expressions for A_{12} and A_{21} matrices:

$$A_{12} = L_{11} \cdot U_{12} \quad (1.6)$$

$$A_{21} = L_{21} \cdot U_{11} \quad (1.7)$$

To compute 1.6 we only have to perform two triangular solves since matrices L_{11} and U_{11} are known. These two tasks are concurrent and they can be computed in parallel. Additionally we can notice the triangular solve routine can be computed in parallel too. The next step is to perform \hat{A}_{22} matrix update and block structure re-ordering:

$$\hat{A}_{22} = A_{22} - L_{21} \cdot U_{12} \quad (1.8)$$





Figure 1.2.: Right-looking parallel LU decomposition

It is clear the algorithm is purely sequential at the first steps when we compute small LU decomposition. Therefore, it can have significant effect on algorithm strong scaling behavior. It should be mentioned that all three parallel implementations have the same problem i.e. they have an inherently sequential part at the beginning of a step.

Figure 1.3 shows results of strong scaling of dense LU factorization performed for a couple of matrices filled with random numbers with different sizes: namely: 5000×5000 , 10000×10000 and 15000×15000 . LAPACK and OpenBLAS libraries were used for the test. One can easily notice that performance of dense LU decomposition

quickly deteriorates with reduction of the problem size. Additional factor that affects performance is strong scaling behavior of the triangular solve that we will discuss in section 1.4.



place for
figure 1.3

Figure 1.3.: Strong scaling of right-looking LU decomposition using LAPACK and OpenBLAS

Both left, right and Crout parallel matrix factorizations have been efficiently implemented in LAPACK (for shared-memory machines) and ScaLAPACK (distributed-memory machines) libraries. Both libraries belong to the Netlib project which is a repository of numerous scientific computing software maintained by AT&T Bell Laboratories, the University of Tennessee, Oak Ridge National Laboratory and other scientific communities [06]. The libraries are built on top of Basic Linear Algebra Subprograms (BLAS) library. Figure 1.4 shows how these three libraries are coupled together.

add
strong
scal-
ing for
ScaLA-
PACK

It is worth noting that BLAS can be considered as a foundation of LAPACK and ScaLAPACK libraries and, thus, it is the primary source of performance improvement. In particular we have to consider DGEMM, DTRSM BLAS subroutines performance because they, together with unblocked LU factorization, compose the core of parallel LU factorization algorithms as we discussed earlier.

place for
figure 1.4

There exist special-purpose, hardware-specific BLAS implementations developed by the hardware vendors i.e. IBM, Cray, Intel, AMD as well as open-source tuned imple-

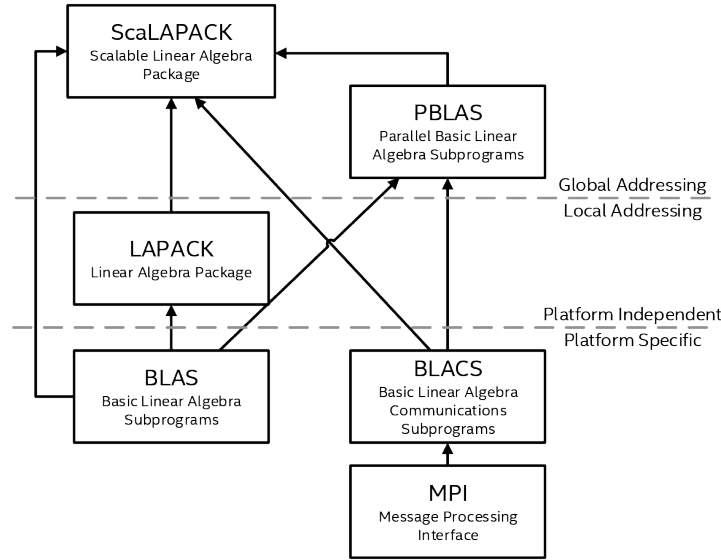


Figure 1.4.: A general view on BLAS, LAPACK and ScaLAPACK libraries [17a]

mentations such as ATLAS, OpenBLAS, etc. We will come back to that discussion later and pay our close attention to a specific choice of a tuned BLAS library in subsection 1.11.

In spite of all advantages of the direct dense solvers i.e. numerical stability and high ratio of floating point operations per memory access, we cannot consider this group of methods as a solver for time integration due to high complexity and storage costs.

1.4. Iterative methods

Iterative methods, especially Krylov subspace methods that we are going to discuss in this section, are well known for their relatively low storage requirements $O(nnz)$ and computation cost $O(N^2)$ in case of sparse linear systems of equations and good condition number. It turns out that sometimes it might be only one way to solve huge systems with millions unknowns.

The most well known methods are Conjugate Gradient (CG) for symmetric positive definite matrices, Minimal Residual Method (MINRES) for symmetric indefinite systems, Generalized Minimal Residual Method (GMRES) for non-symmetric systems of linear equations as well as different variants of GMRES such Biconjugate Gradient

Method (BiCG), Biconjugate Gradient Stabilized Method (BiCGSTAB) and so on.

All Krylov methods solve a system of equation as a minimization problem. For example, the goal of CG algorithm is to minimize the energy functional $f(x) = 0.5x^T Ax - b^T x + c$, whereas, MINRES and GMRES tries to minimize residual norm r_j for x_j in a subspace.

The methods construct an approximate solution of a system as a linear combination of vectors b, Ab, A^2b, A^3b and so on which defines the Krylov subspace. At each iteration we expand the subspace adding and evaluating a next vector in the combination.

Let's consider GMRES, as the most popular and general iterative solver, without preconditioning to just analyze its strong scaling behavior and potential problems.

As we mentioned above GMRES minimizes the residual norm in a subspace U_m .

$$\min_{x \in U_m} \|Ax - b\|^2 \quad (1.9)$$

We can consider a solution vector x in the subspace U_m in a form $x = U_m y$. Thus, equation 1.9 can be written as following:

$$\min_{x \in U_m} \|AU_m y - b\|^2 \quad (1.10)$$

The most natural way to choose a proper subspace U_m is the corresponding Krylov subspace \mathcal{K}_m because it can be easily generated on the fly. However, decomposition of vector x in that subspace can be a problem. Since the subspace \mathcal{K}_m is spanned by the sequence of $b, Ab, A^2b, \dots, A^{m-1}b$ and due to round-off error the sequence can become linear dependent. Therefore, we have to compute and use the orthonormal base of the given Krylov subspace. Saad and Schultz in their work [SS86] used Arnoldi process for constructing an l_2 -orthogonal basis. As the results equation 1.10 can be written in the following form:

$$\min_{x \in U_m} \|U_{m+1} H_{m+1,m} y - \|b\|u_1\|^2 = \min_{x \in U_m} \|H_{m+1,m} y - \|b\|e_1\|^2 \quad (1.11)$$

where H_m is an upper Hessenberg matrix. We can apply Givens rotation algorithm to compute QR decomposition to convert H_m to a strictly upper triangular matrix. Thus,

$$\min_{x \in \mathcal{K}_m} \|Ax - b\|^2 = \min_{x \in U_m} \|Q^T R y - \|b\|e_1\|^2 = \min_{x \in U_m} \left\| \begin{pmatrix} R_m \\ 0 \end{pmatrix} y - \begin{pmatrix} \tilde{b}_m \\ b_{n-m} \end{pmatrix} \right\|^2 \quad (1.12)$$


Given 1.12, we can compute the solution as following:

$$R_m y = \tilde{b}_m \quad (1.13)$$

$$x_m = U_m y \quad (1.14)$$

Because of large computational and storage costs, in case of evaluation of the full Krylov subspace, only small a subspace is computed, typically first 20 - 50 column vectors. Then the algorithm is restarted using the computed approximate solution as a initial guess for the next iteration.

We can see that some operations, for example 1.14, can be efficiently done in parallel. However, operations like sparse triangular solve 1.13 can introduce some effect on strong scaling behavior. Figure 1.5 shows strong scaling performance results of a sparse parallel triangular solver with a two dimensional matrix distribution. Performance considerations of the solver can be found in [Jos+98].



place for
figure 1.5

Figure 1.5.: Performance of sparse triangular solve [Jos+98]

It is interesting to notice that performance of the triangular solver depends on a matrix sparsity structure as well as the matrix size.

Triangular solve 1.13 can be computed in a single processor because matrix R_m is usually small and depends on the number iterations before the restart. In this case the triangular solve can become a bottleneck again.

Figure 1.6 shows strong scaling performance results of the default GMRES solver from the PETSc library. The solver was set up without any preconditioner and 50 iterations as the restart. Additionally no stop criteria was specified except the maximum number iterations which was equal to 100. The *spread* process pinning strategy, described in section 1.2, was used. It is well-known that all iterative methods are memory bound due to indirect memory addressing caused by sparse matrix storage schemes. Hence equal process distribution can help reduce the load on memory channels since it is an obvious bottleneck for this type of applications.

Four matrices were chosen from the GRS matrix set for the tests, namely: cube-64, cube-645, k3-2 and k3-18. The information about the matrices is summarized in table 1.1. As we expected, we can observe strong deviation of our results from the ideal speed-up line when the number of processes exceeds 10.

It should be mentioned that parallelization overheads, introduced by such MPI operations as MPI_Send, MPI_Recv, MPI_Allreduce, etc., also have their impact on performance of the algorithm.

place for
figure 1.6

Other Krylov methods such as CG, for example, scales much better than GMRES. Because of the nature of the CG algorithm the next search direction can be found using a recurrent expression and the algorithm boils down to simple operations such as dot products and matrix vector multiplications. These operations can be easily parallelized and drop of performance comes only from MPI overheads. A quite comprehensive study about parallel CG algorithm performance can be found in [CCV18]. The authors also introduced a deeply pipelined version of CG algorithm that scales even better due to overlapping the time-consuming global communication phase, induced by parallel dot product computations, with useful independent computations [CCV18].

The most important criteria of Krylov methods is convergence rate. The convergence rate of iterative methods strongly depends on a matrix and, in particular, on its condition number. For instance, equation 1.15 shows dependence of the convergence rate from the matrix condition number. It can be clearly seen that a big condition number leads to very slow error reduction and, as the results, to huge number of iterations.

place for
figure 1.6




Figure 1.6.: GMRES strong scaling speed-up

$$\|e^i\|_A \leq 2\left(\frac{\sqrt{k}-1}{\sqrt{k}+1}\right)^i \|e^0\|_A \quad (1.15)$$

where $k = \frac{\lambda_{max}}{\lambda_{min}}$ - condition number of the corresponding matrix.

An obvious solution of such a problem is to reduce the condition number of the original system 1.1. A general method is to transform the original system in such a way that the conditional number of the transformed system gets significantly smaller. The transformation of 1.1 can be done from the left side 1.16 or from the right one 1.17.

$$PAx = Pb \quad (1.16)$$

$$AP(P^{-1}x) = b \quad (1.17)$$

where matrix P is called preconditioner.

As an extreme example we can consider the inverse matrix A^{-1} as the best preconditioner since it directly leads to the solution of the problem 1.1 and, thus, it requires only one iteration. However, it is obvious that computation of inverse A^{-1} is extremely expensive operation and it is not an objective of any iterative methods. That example helps understand and set requirements for preconditioners, namely:

1. cheap to compute e.g. a 5-10 iterations of the corresponding Krylov solver
2. should lead to a small conditioner number of the transposed system
3. should be sparse, otherwise storage requirements will considerably increase

There exist numerous techniques to compute preconditioners given a matrix A e.g. (point) Jacobi, Block-Jacobi, incomplete LU decomposition (ILU), multilevel ILU (ILU(p)), threshold ILU (ILUT), incomplete Cholesky factorization (IC), sparse approximate inverse (SPAI), multigrid as a preconditioner, etc. Almost all methods listed above have some tuning parameters which allow to get a better preconditioner i.e. a smaller condition number of the transformed system. However, it usually leads to increase of computational and storage costs.

Some methods can works particularly well for matrices derived from certain PDEs e.g. Poisson, NavierStokes, etc. problems discretized using the cartesian grid. However sometimes it can take a considerable amount of time to choice right parameters for a certain preconditioning algorithm. It can become a challenge to fulfill all requirements 1, 2, 3 mentioned above.

Table

Table [] shows results of different preconditioning algorithms application to our test case. It can be seen that some algorithms failed even after tuning.

It interesting to notice that Gupta, Koric, and George came to approximately the same results working on their set of matrices in their work [GKG09]. They observed that preconditioned iterative solvers worked efficiently only for 2 out 5 cases in contrast to direct sparse solvers.

We can summarize that it is vital to perform careful parameter tuning of any preconditioning algorithms combining results from [table] and [GKG09]. In general the search can take a considerable amount of time. Moreover, it becomes impractical for time integration problems where topology of an underlying problem and, as the results, the computational mesh, discretization, Jacobian matrix can be changed over time of a simulation. It is obvious that parameters chosen for a particular time step can become not optimal for consecutive steps and, at the end, it can lead to divergence. If divergence happens at any time step the entire time integration algorithm fails and the simulation has to be restarted with different preconditioning parameters or with a different preconditioning algorithm.

Table with comparisons of different preconditioning for our test cases

By and large we come to a conclusion that preconditioned iterative solvers are not robust and thus cannot fully fulfill requirements listed in section 1.

1.5. Direct sparse methods

Direct sparse methods combine main advantages of direct and iterative methods i.e. numerical robustness and usage of sparsity structures. As a results, there is no need for preconditioning and the computation complexity is $O(n^2)$ [Wu12]. The problem is that storage cost can significantly increase during factorization i.e. the inverse of a sparse matrix can be sufficiently dense. To reduce storage space of LU decomposition this group of methods performs fill-in reduction reordering as a pre-processing step before actual factorization. If storage space is still huge even after fill-in reduction reordering out-of core factorization can be used where partial results are stored in the secondary memory.

The most widely known sparse direct method is multifrontal method introduced by Duff and Reid in their work [DR83]. Multifrontal method is an improved extension of a frontal method [IB70] that can compute independent fronts in parallel. A front, or also called frontal matrix, can be considered as small dense matrix which is a result of Gaussian Elimination for a particular column. The algorithm, in fact, is as a variant of Gaussian Elimination process. There also exist left-looking and right-looking sparse direct methods. The difference between all of them is explained and can be found in [PT01].

In order to understand and analyze strong scaling behavior of the algorithm we have to briefly discuss the theory of the method. For simplicity we will assume that matrix A is real symmetric and LU decomposition boils down to the Cholesky factorization 1.18. It allows us to focus on the Cholesky factor L and its sparsity pattern only.

$$A = LDL^T \quad (1.18)$$

The algorithm usually starts with symbolic factorization to predict sparsity pattern of L . Once it is done the corresponding elimination tree has to be constructed.

Figure 1.7 shows an illustrative example of a sparse matrix and its Cholesky factor from [Liu92]. The solid circles represent original non-zero elements whereas hollow

place for
figure 1.7

$$A = \begin{pmatrix} 1 & a & & & & & & & \\ 2 & & b & & \bullet & & & & \\ 3 & & & c & & \bullet & & & \\ 4 & & \bullet & & d & & & & \\ 5 & & & \bullet & & e & & \bullet & \\ 6 & & \bullet & & & & f & & \\ 7 & \bullet & & & & & & g & \bullet \\ 8 & \bullet & & \bullet & \bullet & & & & h \\ 9 & \bullet & & & \bullet & \bullet & & & i \end{pmatrix} \quad L = \begin{pmatrix} 1 & a & & & & & & & \\ 2 & & b & & & & & & \\ 3 & & & c & & & & & \\ 4 & & \bullet & & d & & & & \\ 5 & & & \bullet & & e & & & \\ 6 & & \bullet & & & & \circ & \bullet & f \\ 7 & \bullet & & & & & & & g \\ 8 & \bullet & & \bullet & \bullet & \bullet & & \circ & \bullet & h \\ 9 & \bullet & & & \bullet & & \bullet & \bullet & \circ & i \end{pmatrix}$$

Figure 1.7.: An example of a sparse matrix and its Cholesky factor [Liu92]

ones define fill-in factors of L .

The elimination tree is a crucial part of the method. It can be considered as a structure of n nodes that node p is the parent of j if and only if it satisfies equation 1.19. It is worth pointing out the definition 1.19 is not only one possible and one can define a structure of an elimination tree in a different way as well. As an example one can find a definition of a general assembly tree in [Liu92] proposed by Liu.

$$p = \min(i > j | l_{ij} \neq 0) \quad (1.19)$$

It is important to notice that node p represents elimination process of the corresponding column p of matrix A as well as all dependencies of column p factorization on the results of its descendants.

Given definition 1.19 we can build the corresponding elimination tree as it is shown in figure 1.8.

The fundamental idea of multifrontal method spins around frontal and update matrices. A frontal matrix is used to perform Gaussian Elimination for a specific column j . It is a sum of a frame and update matrices as it can be seen from equation 1.20

place for figure 1.8




Figure 1.8.: The elimination tree for the matrix example in Figure 1.7 [Liu92]

$$F_j = Fr_j + \hat{U}_j = \begin{bmatrix} a_{j,j} & a_{j,i_1} & a_{j,i_2} & \dots & a_{j,i_r} \\ a_{i_1,j} & & & & \\ a_{i_2,j} & & & & \\ \vdots & & 0 & & \\ a_{i_r,j} & & & & \end{bmatrix} + \hat{U}_j \quad (1.20)$$

where i_0, i_1, \dots, i_r are the row subscripts of non-zeros in L_{*j} with $i_0 = j$ and r is number of off-diagonal non-zero elements.

The frame matrix Fr_j is filled with zeros except the first row and column. The first row and column contain non-zeros elements of the j th row and column of the original matrix A . Because we consider matrix A to be symmetric the frame matrix is square and symmetric as well.

In order to describe parts of the elimination tree we will use the notation $T[j]$ to represent all descendants of the node j in the tree and node j itself. In this way we can define the update matrix \hat{U}_j as following:

$$\hat{U}_j = - \sum_{k \in T[j]-j} \begin{bmatrix} l_{j,k} \\ l_{i_1,k} \\ \vdots \\ l_{i_r,k} \end{bmatrix} [l_{j,k} \quad l_{i_1,k} \quad \dots \quad l_{i_r,k}] \quad (1.21)$$

The update matrix \hat{U}_j is, in fact, can be considered as the second term of the Schur complement i.e. update contributions from already factorized columns of A .

The subscript k represents descendant columns of node j . Thus we include and consider only those elements of descendant columns which correspond to the non-zero pattern of the j th column that we are currently factorizing.

Let's consider the partial factorization of 2-by-2 block dense matrix to better understand essence of update matrix \hat{U}_j .

$$A = \begin{bmatrix} B & V^T \\ V & C \end{bmatrix} = \begin{bmatrix} L_B & 0 \\ VL_B^{-T} & I \end{bmatrix} \begin{bmatrix} I & 0 \\ 0 & C - VB^{-1}V^T \end{bmatrix} \begin{bmatrix} L_B^T & L_B^{-1}V^T \\ 0 & I \end{bmatrix} \quad (1.22)$$

Again we assume that B has already been factorized and can be expressed as:

$$B = L_B L_B^T \quad (1.23)$$

The Schur complement from equation 1.22 can be viewed as the original sub-matrix C and update $-VB^{-1}V^T$. It can be written in a vector form as well:

$$-VB^{-1}V^T = -(VL_B^{-T})(L_B^{-1}V^T) = -\sum_{k=1}^{j-1} \begin{bmatrix} l_{j,k} \\ \vdots \\ l_{n,k} \end{bmatrix} [l_{j,k} \quad \dots \quad l_{n,k}] \quad (1.24)$$

As it can be easily seen that equations 1.24 and 1.21 are identical. The difference is that equation 1.21 exploits sparsity of the corresponding row and column of L and thus masks unnecessary information.

We can also notice from equation 1.22 that the frame matrix Fr_j corresponds to the block matrix C and brings information from the original matrix A whereas matrix \hat{U}_j adds information about the columns that have already been factorized.

As soon as the frontal matrix F_j is assembled i.e. we have the complete update of column j , we can perform elimination of the first column and get non-zero entries of factor column L_{*j} .


Let's denote \hat{F}_j as a result of the first column factorization of the frontal matrix F_j . Then we can express the results as following:

$$\hat{F}_j = \begin{bmatrix} l_{j,j} & \dots & 0 \\ \vdots & I & \\ l_{i_{r,j}} & & \end{bmatrix} \begin{bmatrix} 1 & \dots & 0 \\ \vdots & U_j & \\ 0 & & \end{bmatrix} \begin{bmatrix} l_{j,j} & \dots & l_{i_{r,j}} \\ \vdots & I & \\ 0 & & \end{bmatrix} \quad (1.25)$$

where sub-matrix U_j represents the full update from all descendants of node j and node j itself. Equation 1.26 express the sub-matrix U_j in a vector form.

$$\hat{U}_j = - \sum_{k \in T[j]} \begin{bmatrix} l_{i_1,k} \\ \vdots \\ l_{i_1,k} \end{bmatrix} [l_{i_1,k} \quad \dots \quad l_{i_1,k}] \quad (1.26)$$

Together with the frontal F_j and update \hat{U}_j matrices, the update column matrix U_j (also called contribution matrices) forms the key concepts of the multifrontal method. To consider the importance of sub-matrix U_j let's consider an example illustrated in Figure 1.9.



place for
figure 1.9

Figure 1.9.: Information flow of the multifrontal method

We assume that factorization of columns A and B have already been done and corresponding contribution matrices U_A and U_B have been computed. From equation 1.26 we have already known that both U_A and U_B contain the full updates of all their descendants including updates from factorization of columns A and B as well. Therefore update column matrices U_A and U_B have already got all necessary information to construct update matrix \hat{U}_C . The detailed proof and careful explanation can be found in [Liu92].

It might happen that we do not need all rows and columns of U_A and U_B i.e. we need only some subset of them, because of sparsity of column C. It is also important

to place all necessary rows and columns of matrices U_A and U_B in a right place within matrix \hat{U}_C . For that reason, an additional matrix operation, called *extend-add*, must be introduced.

Let's consider an example from [Liu92] of an extend-add operation for 2-by-2 matrices R and S which correspond to the indices 5,8 and 5,9 of some matrix B , respectively.

$$R = \begin{bmatrix} p & q \\ u & v \end{bmatrix}, S = \begin{bmatrix} w & x \\ y & z \end{bmatrix} \quad (1.27)$$

The result of the operation is going to be a 3-by-3 K matrix which looks as following:

$$K = R \oplus S = \begin{bmatrix} p & q & 0 \\ u & v & 0 \\ 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} w & 0 & x \\ 0 & 0 & 0 \\ y & 0 & z \end{bmatrix} = \begin{bmatrix} p+w & q & x \\ u & v & 0 \\ y & 0 & z \end{bmatrix} \quad (1.28)$$

Hence we can express formation of the frontal matrix F_j using the extend-add operation and all direct children of node j in the following way:

$$F_j = \begin{bmatrix} a_{j,j} & a_{j,i_1} & a_{j,i_2} & \dots & a_{j,i_r} \\ a_{i_1,j} & & & & \\ a_{i_1,j} & & & & \\ \vdots & & & 0 & \\ a_{i_r,j} & & & & \end{bmatrix} \oplus U_{c_1} \oplus \dots \oplus U_{c_s} \quad (1.29)$$

where c_1, c_2, \dots, c_n are indices of direct children of the node j .

Now it can be clearly seen that the resultant frontal matrix F_j is a small dense one and it can be efficiently computed using BLAS level 3 subroutines.

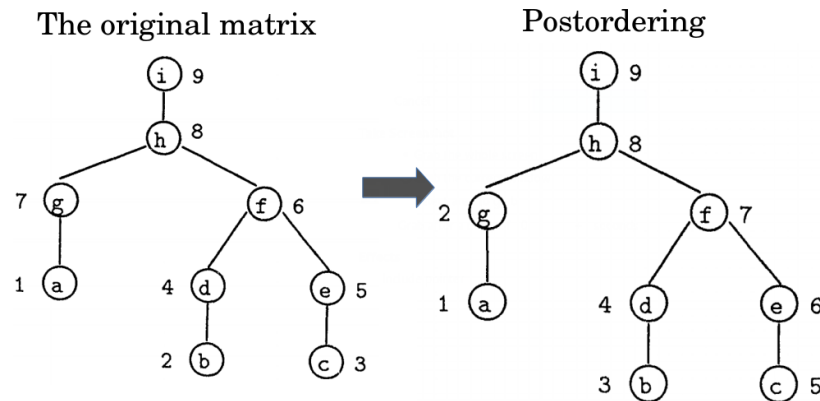
After factorization we have to build the contribution matrix U_j i.e. add columns and rows of $U_{c_1}, U_{c_2}, \dots, U_{c_s}$ to U_j that have not been used in factorization of F_j due to sparsity of column j . After that we can continue to move up along the tree. The complete update matrices grow in size as we move to the top of the tree. Therefore they have to be stored in a sparse matrix format to stay within memory constraints of the computer.

Another important aspect is storage and manipulation of frontal and contribution matrices. Sometimes we have to store contribution matrices produced in previous steps into some temporary buffer and efficiently retrieve them later during factorization. This

can require some matrix re-ordering. In case of symmetric matrices, one can apply postordering on a tree to be able to use the stack data structure to alleviate the process of contribution matrix manipulations during factorization. A tree postordering is based on topological ordering and it has been proven that it is equivalent to the original matrix ordering and thus leads to the same filled graph [Liu92]. We refer to the original matrix ordering as the ordering received from fill-in reduction operation.


sentence
refactor-
ing

A tree postordering means that a node is ordered before its parent and, additionally, nodes in each subtree are numbered consecutively. Figure 1.10 shows an example of postordering applied to the elimination tree of the matrix from figure 1.7. The results of this can be seen in figure 1.11 where consecutive *push* and *pop* operations are efficiently used during factorization and thus simplify the program logic.



place for
figure
1.10

Figure 1.10.: An example of matrix postordering from [Liu92]



place for
figure
1.11

Figure 1.11.: The stack contents for the postordering [Liu92]

We can see the algorithm requires to perform some preprocessing steps in order to estimate the size of working space for matrix manipulations. If the working space

has not been predicted correctly the algorithm will terminate during factorization. Additionally it can happen that even with the correct estimation we can be run out of space in the main memory, in case of huge sparse matrices. This fact can require to use the secondary memory and, as a result, the execution time will increase significantly. Therefore, different optimal postordering schemes have been proposed which allow to shrink the amount of space needed during factorization [Liu86] [Liu88]. Some schemes, for example elimination tree rotations [Liu88], can lead to deep and unbalanced trees which might have their negative effect on task parallelism as we will see later.

In general, the estimation of working space can be tricky due to pivoting. Because pivoting happens only during the numerical factorization it is not always possible to estimate enough space correctly beforehand. There exist some heuristics which allow to use some numerical matrix information during symbolic factorization to better predict the amount of required space [Gup00].

It can be clearly observed the method consists of three distinct phases, namely: analysis, numerical factorization and solution. The analysis phase includes all pre-processing steps that have been discussed above i.e. fill-in reduction, postordering, symbolic factorization, building elimination tree and so on. During the numerical factorization phase the L and D (or U) parts of a matrix A are computed based on sequence of dense factorization on frontal matrices. At the solution step, the solution vector x is computed by means of backward and forward substitutions (equations 1.2 and 1.3).

In practice, an improved version of multifrontal method, called supernodal method, is used. The idea of the supernodal method is to shrink the final elimination tree by grouping some particular nodes/columns in one computational unit. As a result, more useful floating point operations per memory access can be performed by eliminating few columns at once within the same frontal matrix.

A supernode is formed by a set of contiguous columns with identical off-diagonal sparsity structure forms. Thus, a supernode has few important properties. Firstly, it can be expressed as a set of indices, namely: $\{j, j+1, \dots, j+t\}$, where node $j+k$ is parent of $j+k-1$ in the elimination tree. Secondly, the size of the supernodal frontal matrix is equal to the frontal matrix of the j th column within a supernode. As an example, Figure 1.12 shows a postordered matrix A and its Cholesky factor L as well as the corresponding supernodal elimination tree.

Equation 1.30 expresses the building process of a frontal matrix of a supernode. In contrast to 1.29, the frame matrix \mathcal{F}_j contains more dense rows and columns. As before,

place for
figure
1.12




Figure 1.12.: An example of a supernodal elimination tree [Liu92]

we use *extend-add* operation to get the full block update from children contribution matrices.

It should be mentioned there exist more sophisticated variants of supernodes. Most of the time, it intends to improve efficiency of the algorithm. Liu pointed out that supernodes could be defined without using the contiguous constraints [Liu92]. On another hand, Wu defines supernodes corresponded to separators from the nested dissection step [Wu12] which was used for fill-in reduction.


check grammar

$$\mathcal{F}_j = \begin{bmatrix} a_{j,j} & a_{j,j+1} & \dots & a_{j,j+t} & a_{j,i_1} & \dots & a_{j,i_r} \\ a_{j+1,j} & a_{j+1,j+1} & \dots & a_{j+1,j+t} & a_{j+1,i_1} & \dots & a_{j+1,i_r} \\ \vdots & \vdots & \dots & \vdots & & & \\ a_{j+t,j} & a_{j+t,j+1} & \dots & a_{j+t,j+t} & a_{j+t,i_1} & \dots & a_{j+t,i_r} \\ a_{i_1,j} & a_{i_1,j+1} & \dots & a_{i_1,j+t} & & & \\ \vdots & \vdots & \dots & \vdots & & & 0 \\ a_{i_r,j} & a_{i_r,j+1} & \dots & a_{i_r,j+t} & & & \end{bmatrix} \quad \clubsuit U_{c_1} \quad \clubsuit \dots \quad \clubsuit U_{c_s} \quad (1.30)$$

Up to this point we have already seen all key concepts of the multifrontal method and discussed how the algorithm works. We will move to the discussion of parallelization of the method.

The elimination tree, in fact, represents dependencies among columns. Conversely, the tree also shows independent steps of elimination process. Hence the tree forms independent problems that can be executed in parallel. Task parallelism is the main and primary source the algorithm parallelisation. Figure 1.13 shows task parallelism,

for the example given in Figure 1.12, where each color represents a set concurrent tasks.



place for
figure
1.13

Figure 1.13.: Parallel steps of the multifrontal method based on the example in Figures 1.12

For example, nodes on separate branches of the tree are totally independent and can be processed in parallel. However, as soon as at least two branches run into the same node it forms a dependency and we have to wait all contribution matrices of its children and cannot proceed further.

We can observe the amount of task parallelism is rapidly decreasing while moving towards the root along the tree. Once we reach the root of the tree the algorithm becomes totally sequential. This fact can play the significant role in strong scaling behavior of the method.

We developed two simple models based on perfectly balanced binary trees to better understand strong scaling of the algorithm. The main concept of the models is so-called cost per level or cost per node. This idea is similar to the recursion trees in [Cor+09] which explains and computes complexity of recurrent algorithms.

Figure 1.14a represents the first model where we keep the same cost per level whereas the second model (Figure 1.14b) simulates quadratic cost decay from level to level. Additionally we assume that computational cost distributed uniformly between nodes at the same level for both models.

place for
figure
1.14




Figure 1.14.: Simple parallel models of the multifrontal method

We have to say that our models mimic only numerical factorization and do not include time spent on any per-processing steps, for example, fill-in reduction reordering. A cost per level can be interpreted in different ways e.g. increase of partial factorization time due to growth of frontal matrices in size, time increase spent on numerical pivoting, increase of MPI communication overheads due to growth of contribution matrices, etc. It should be mentioned that real computer implementations of the multifrontal algorithm (MUMPS, SuperLU, etc.) are quite sophisticated in many aspects and our models do not have any intention to analyze performance of a particular package. Instead the objective of these models is to show possible strong scaling behavior and possible bottlenecks.

sentence
refactor-
ing

We will consider only task parallelism at the beginning to a first approximation and later we will discuss how additional data parallelism can affect algorithm performance.

Instead of coloring given in Figure 1.13, we assume that each level has the same color and thus can be executed fully in parallel if we have enough processing elements. We cannot go to the next level till the current one has not been completed yet i.e. free processing elements, that do not have nodes to execute at the current level, have to wait.

As we mentioned above the root of the tree can be processed purely sequentially if we only consider task parallelism. As a first approximation, time spent on the root factorization determines the minimal execution time according to the Amdahl's law [Wik18a]. More precisely, the minimal execution time is equal to a sum of time spent on single node partial factorization at each level. This time determines the asymptote on the corresponding speed-up graph.

We considered a perfectly balanced tree with 16 levels, 65535 nodes and the maxi-

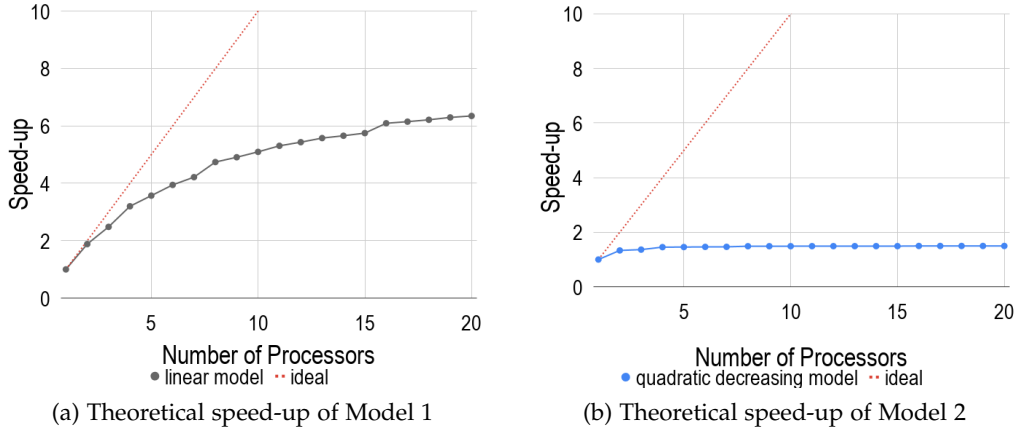


Figure 1.15.: Theoretical speed-up

num of 20 processing elements as an example. The numerical results of linear and quadratic models can be viewed in Figures 1.14a and 1.14b, respectively. The figures show a rapid drop of performance, especially in case of the quadratic model. Table 1.4 demonstrates the maximum potential speed-up, having 32768 processing elements which is equal to the number of leaves of the tree, against the speed-up we have got using only 20 processing elements.

	20 PEs	32768 PEs
Model 1	6.3492	8.0000
Model 2	1.4972	1.5000

place for
figure
1.15

Table 1.4.: Potential speed-up of linear and quadratic models

We can see that model 1 still has some potential to grow whereas the second model has already reached its asymptote and further increase of processing elements does not make sense. In spite of a potential growth of the first model, both models have very low parallel efficiency even with 20 processing elements which can be observed from table 1.5.

Both models shows that computational intensity per node grows from bottom to top. It is easy to conclude from Figure 1.14 that intensity per node is equal $t/2^i$ and $t/2^{2i}$ for the first and second models, respectively (where i is a level of the tree). It reflects

	20 PEs
Model 1	0.3175
Model 2	0.0749

Table 1.5.: Efficiency of linear and quadratic models using 20 PEs

that the most intensive part of the method is centered on the top part of the tree i.e. first few level. Liu discussed application of the multifrontal method to a $k - by - k$ regular model problem with nine-point difference operator in his paper [Liu92]. He observed that factorization of the last 6 nodes took slightly more than 25% of the total amount of arithmetical operations. As a comparison, table 1.6 shows fractions of time spent on processing first few top levels of our models: 1 and 2.

	Model 1	Model 2
Level 0	6.25%	50.00%
Level 1	12.50%	75.00%
Level 2	18.75%	87.50%

Table 1.6.: Distribution workload per level in case model 1 and 2

As we can see, the result of our first model is relatively close to 25% and, therefore, it looks quite optimistic. However, the second model shows that 87% of workload is focused on the top part of the tree and, as a result, we can consider that model as extremely pessimistic.

By and large, reduction of time spent on the top nodes is a way to improve strong scaling behavior. To do so, data parallelism can be additionally exploited for these nodes. It is worth noting that data parallelism at bottom levels does not make sense because it leads to increase of granularity there and thus increase communication overheads which can lead to significant performance drop.

Figure 1.16 shows an example of two types of parallelism applied to the algorithm. First of all, we can see the leaves are grouped in subtrees and a single PE is assigned to each subtree. Other nodes are distributed among three different types. Nodes of the first type uses task parallelism only, which is induced by the tree, and each node is executed in a single processor. The second type exploits data parallelism with 1D block row distribution among the processors. The root belongs to the third type where data parallelism is used with 2D block cyclic distribution. The details of MUMPS parallelism

management is carefully explained and can be found in [OA07].




Figure 1.16.: MUMPS parallelism management in case of 4 PEs [OA07]

All the techniques mentioned above were designed to improve strong scaling behavior by splitting the most intensive parts among all available processors. Going back to our models, we can also think about that in a slightly different way, namely: *data parallelism helps to re-distribute cost per node/level on the corresponding elimination tree*. However, we have to notice that efficiency of data parallelism totally depends on sizes of frontal matrices at the top part of the tree. In case of skinny sparse matrices, oversubscription of processing elements can lead to strong performance penalties as we could see from section 1.3. A machine-dependent minimal frontal matrix size was introduced in MUMPS in order to control whether to use ScaLAPACK at the root node or not [17b]. It can happen that the algorithm uses only task parallelism, due to the threshold, and, as a results, scaling will only depend on the tree structure that can be deep and unbalanced.

Figure 1.17 shows comparison of strong scaling between model 1 and parallel numerical factorization of the matrix *cant* (Table 1.2) done with using MUMPS library. The sparsity pattern before and after fill-in reduction is shown in figure 1.18.

place for
figure
1.16

add some
examples
to the
appendix

place
for fig-
ure 1.17.
Show
standard
deviation

place for
figure
1.18

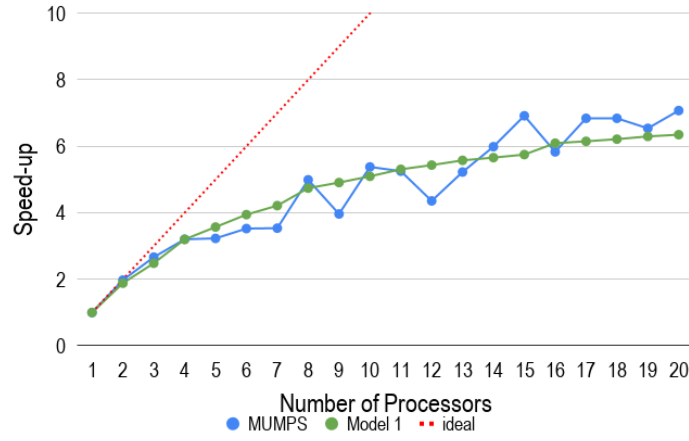


Figure 1.17.: Comparison between model 1 and numerical factorization of the matrix *cant* using MUMPS library

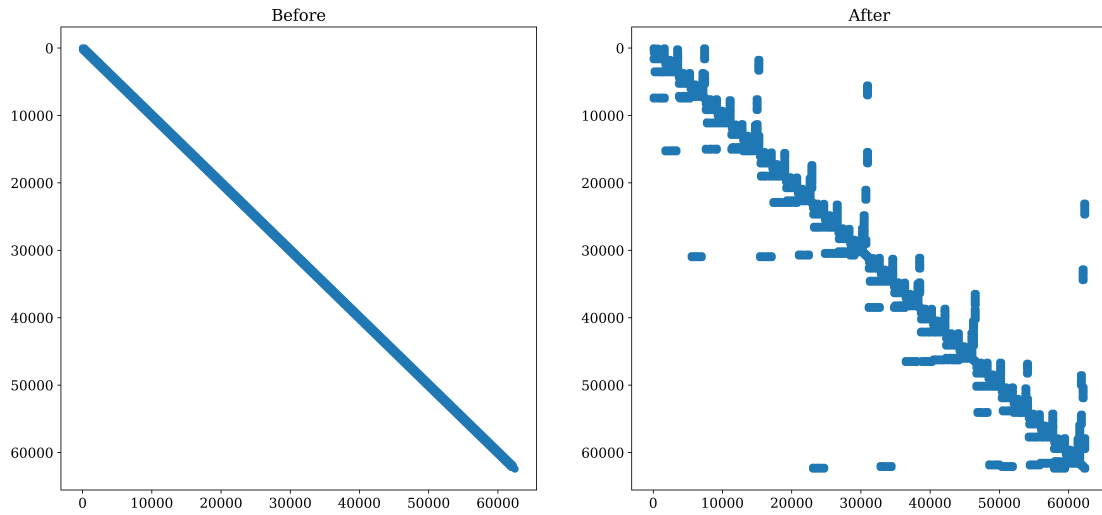


Figure 1.18.: Sparsity structure of the matrix *cant* before and after fill-in reduction

As we can see, our model and the experiment show the same trend and the results are pretty much close to each other. However, our model takes into consideration only task parallelism whereas MUMPS exploits both data and task parallelism. Additionally, we have to mention that our model 1 works well only for relatively big sparse matrices. It can be quite inaccurate in case of small/skinny sparse systems.

In general, it is possible to refine our models and make them more accurate, using Bulk Synchronous Parallel (BSP) approach, for example. However, it will require to possess a real postordered elimination tree, extracted from a specific implementation of multifrontal method, together with information about supernodes sizes. We strongly believe the new enhanced model can explain the jagged strong scaling behavior of the MUMPS solver that we can observe in figure 1.17. But, this approach seems to be quite cumbersome and requires to delve into the source code of a particular library. It is needless to say that data can be retrieved only during run time and only after the analysis phase. This makes it less valuable and we can see that it is definitely a wrong way to go.

There are few important aspects to discuss at the end of the section. Numerical robustness is the main advantage of the multifrontal method. It does not require any preconditioner to solve a system of equations. As we discussed in the previous section, tuning a specific preconditioning algorithm can take a considerable amount of time, especially in case of our systems. As another advantage, the method (heavily) exploits matrix sparsity which lowers computational complexity up to $O(n^2)$. In case of massively huge matrices, the algorithm can utilize the secondary memory which sometimes is only one way to solve a system.

We can conclude, from the analysis above, the method has inherently bad scaling behavior and it is quite sensitive to a matrix structure. We will see later that it is almost impossible to predict the saturation point i.e. a point after which performance either drops or stays at the same level. We assume that scaling becomes better with growth of a matrix size. However, we cannot expect such behavior for small and medium systems.

Secondly, we can see the algorithm requires many pre-processing steps to be done before numerical factorization phase. All these steps must run in parallel and be highly scalable. Apart from performance constraints of the steps, they must lead to wide and well balanced elimination trees which becomes crucial during the numerical phase.

Lastly, the algorithm can fail due to incorrect working space prediction. As a results, factorization has to be restarted with some modification of input solver parameters.

1.6. Hybrid method

We have observed almost all available methods and we could see that none of them can fully cover all our requirements at once, namely:

- robustness
- numerical stability
- parallel efficiency
- open source licenses

The analysis from sections 1.4 and 1.4 shows that iterative methods scale much better in contrast to sparse direct ones. However, they are only efficient in case of very well preconditioned systems. We showed in section 1.4 that search of preconditioning parameters usually takes lots of time and efforts. Additionally, we cannot guarantee that the settings found for our GRS matrix set will always work well in subsequent steps of time integration or for other different simulations.

Sparse direct methods do not have such a problem. They always produce the right solution. The methods can only fail in case of underestimation of the working space due to numerical pivoting during the numerical factorization phase. In order to cope with that problems, some implementations of direct sparse methods provide two options to the user, namely: to increase predicted working space by some factor e.g. 2, 3, 4, etc. or to lower constraints of numerical pivoting which allows small numerical values to stay on the diagonal.

The drawback of the second option is that it can lead to out-of-core execution with using the secondary memory which makes numerical factorization significantly slow. While the second option has lower chance of out-of-core factorization it can lead to a numerically inaccurate solution.

After many considerations we decided to stick to the sparse direct solvers because **robustness** criteria had the highest priority in our case. To circumvent problems mentioned above, we proposed a so-called hybrid solver, in spite of the fact that the definition of *hybrid linear solvers* had already been used in scientific computing literature in a slightly different way [RBH12]. *The idea is to switch off numerical pivoting (or*

significantly lower the constraints) of sparse direct solvers and use the resultant LU decomposition as a preconditioner for an iterative method, for example GMRES.

```
1 # compute LU decomposition with a sparse direct solver
2 LU = SparseDirectSolver(matrix=A, pivoting="switch_off")
3
4 # compute inverse of A using backward-forward substitutions
5 # solve: LU * A_inv = I
6 A_inv = ComputeInverse(decomposition=LU)
7
8 # apply a Krylov method to a preconditioned system
9 # i.e A_inv * A * x = A_inv * b
10 GMRES(matrix=A, rhs=b, preconditioner=A_inv)
```

Listing 1.1: A pseudo-code of the Hybrid approach

According to our primary tests, the hybrid approach showed us that it required from 1 to 5 iterations of the GMRES solver on average to converge to a desired residual.

write
which
solver we
used

The main problem of our approach is parallel efficiency because sparse LU decomposition takes the most of computational time. We discussed reasons of possible bad strong scaling behavior of sparse direct solvers in section 1.5. We could see, in case of the multifrontal method, these methods consist of multiple steps and implementation of each step has its strong effect on parallel performance. We also mentioned that the main source of performance improvement is data parallelism and it can be achieved in many different ways. Hence, performance of the same method can vary from library to library.

In the next section we are going to investigate all available open-source implementations of sparse direct solvers, compare their efficiency and choose one of them. At the beginning, we will only consider libraries that have their direct interface to PETSc [Bal+18]. PETSc is a scientific numerical library that contains various algorithms and methods, especially the Krylov methods. It is highly efficient in parallel and provides numerous interfaces to other libraries such as MUMPS, SuperLU, Hypre, PaStiX, ViennaCL and etc.

The subsequent sections will be dedicated to tuning and optimization of a specific library with the aim to reduce execution time.

proofreading

1.7. Choice of a Sparse Direct Solver Library

Fair to say, there is no single algorithm or software that is best for all types of linear systems [Li18]. Nowadays there exist many different sparse direct solvers on the market. Some of them are tuned for specific linear systems i.e. symmetric positive definite, systems with symmetric sparsity pattern, system with complex numbers, etc., some are targeted for the most general cases. Some packages can handle data parallelism in different ways even within the same library depending on the system size and other criteria. Hence, parallel performance highly depends on a specific implementation of a method. Table 1.7 displays a short summary of almost all available and well-known packages, at the time of writing, in this field based on works [Li18] and [Bal+18].

Package	Method	Matrix Types	PETSc Interface	License
Clique	Multifrontal	Symmetric	Not Officially	Open
MF2	Multifrontal	Symmetric pattern	No	-
DSCPACK	Multifrontal	SPD	No	Open*
MUMPS	Multifrontal	General	Yes	Open
PaStiX	Left looking	General	Yes	Open
PSPASES	Multifrontal	SPD	No	Open*
SPOOLES	Left-looking	Symmetric pattern	No	Open*
SuperLU_DIST	Right-looking	General	Yes	Open
symPACK	Left-Right looking	SPD	No	Open
S+	Right-lookin	General	No	-
PARDISO	Multifrontal	General	No	Commercial
WSMP	Multifrontal	General	No	Commercial

add foot-
notes of
Open*

Table 1.7.: List of packages to solve sparse linear systems using direct methods on distributed memory parallel machines [Li18], [Bal+18]

We only listed libraries that can run on distributed memory parallel machines. Almost all of them also support shared memory environment in some degree. Nonetheless there also exist libraries that run either only sequentially (UMFPACK, SPARSE, TAUCS, SuperLU) or only on shared memory machines (PanELLT, SuperLU_MT).

We can see, from table 1.7, that only MUMPS, PaStiX and SuperLU_DIST cover all our initial requirements: open-source license and direct interface to the PETSc library. However, these libraries implement different sparse direct methods, namely: multifrontal, left-looking and right-looking, respectively. Moreover, they handle partial pivoting in different ways.

It is known that partial pivoting is necessary to achieve a good numerical accuracy during Gaussian Elimination. It interchanges rows and columns of a matrix in such a way to avoid small numerical values along the diagonal. In case of sparse direct solvers, the numerical pivoting, in run-time, usually distorts all predictions that have been made during the analysis phase and can lead to significant fill-in and load unbalanced, with respect to floating-point operations, during factorization. Hence, implementation of numerical pivoting, especially during parallel execution, plays the most important role in performance for this group of methods.

Both PaStiX and SuperLU_DIST libraries use so-called static pivoting where the pivot order is chosen before numerical factorization and kept fixed during factorization.

The main advantage of static pivoting is that it allows to better optimize the data layout, load balance, and communication scheduling [11]. However, it leads to a higher risk of numeric instability. Therefore, both PaStiX and SuperLU_DIST provide a few ways to perform the solution refinement.

For instance, SuperLU_DIST uses diagonal scaling, setting very tiny pivots to larger values, and iterative refinement (listing 1.3). While PaStiX allows the user to choose a refinement strategy between GMRES, CG (for SPD systems) and iterative refinement as well. At this point it is interesting to notice that we came to the same conclusion as the PaStiX developers with respect to the solution refinement using Krylov iterative methods.

Iterative refinement, shown in listing 1.3, is claimed to converge to a decent precision within 2 or 3 steps in work [ADD89]. However, in practice, we noticed that the iterative refinement can work not as expected, especially in case of lowered partial pivoting constrains.

can we
use de-
cent here?

For completeness, we have to mention that the variable *too_large* is, in fact, an estimation of the backward error [ADD89] which can be expressed as following:

$$\frac{|b - A\hat{x}|_i}{(|b| + |A||\hat{x}|)_i} \quad (1.31)$$

where \hat{x} is the computed solution and $|\bullet|$ is the element-wise module operation.

```

1 # perform analysis and numerical factorization
2 # phases
3 LU = SparseDirectSolver(matrix=A)
4
5 # compute initial solution
6 x = Solve(factorization=LU, rhs=b)
7
8 # compute initial residual
9 r = A * x - b
10
11 while r > too_large:
12     # find correction
13     d = Solve(factorization=LU, rhs=r)
14
15     # update solution
16     x = x - d
17
18     # update residual
19     r = A * x - b

```

Listing 1.2: A simple iterative refinement

In contrast to PaStiX and SuperLU_DIST, MUMPS performs partial pivoting in run-time during the numerical factorization phase. To limit the amount of numerical pivoting, and stick better to the sparsity predictions done during the symbolic factorization, partial pivoting can be relaxed, leading to the partial threshold pivoting strategy [17b].

listing of
iterative
refine-
ment

A pivot $|a_{i,i}|$ is accepted if it satisfies:

$$|a_{i,i}| \geq u \times \max_{k=i \dots n} |a_{k,j}| \quad (1.32)$$

where u is value between 0 and 1.

To improve solution accuracy, MUMPS, as PaStiX and SuperLU_DIST, provides the iterative refinement as a post-processing step as well.

The most important feature which MUMPS introduces is so-called delayed pivots. It can happen that equation 1.32 cannot be satisfied within a fully-summed block of a frontal matrix (equation 1.30) and we also cannot consider elements outside the block since the corresponding rows are not fully-summed. In this case, some rows and columns will remain unfactored, or delayed, in the front. They are going to be sent the frontal matrix of the parent, as part of the contribution block and the process will repeat. The delayed pivot approach helps to improve numerical accuracy, however, it causes additional fill-in in the parent node.

In spite of obvious complexity of dynamic partial pivoting, MUMPS allows the user to explicitly control run-time behavior of the algorithm due to partial threshold pivoting strategy. This provides an opportunity for optimization and tuning in some degree.

PETSc (version 3.10) provides the full interface to both SuperLU_DIST and MUMPS, whereas the interface to the PasTiX library is quite limited. In fact, in case of PasTiX, the user can only control the number of threads per MPI process and a level of verbosity, which makes this library to be less interesting for our subsequent (following) research.

data parallelism difference between MUMPS and SuperLU

In order to evaluate the overall parallel performance of the libraries, we performed a few flat-MPI tests with the GRS matrix using the HW1 machine. Before testing, we downloaded and configured the libraries, MUMPS version 5.1.2, PasTiX version 6.0.0, SuperLU_DIST version 5.4, within the PETSc environment with their **default settings**. As a profiling tool, we used the internal PETSc profiler. A time limit of 15 minutes was set up for each test case to prevent blocking of a compute node in case if out-of-core execution. Results of the tests are summarized in tables 1.8, 1.9, 1.10 and in appendix A. Numerical values in tables are given in **seconds**.

We ran into a few problems with the SuperLU_DIST library during the tests. Firstly, factorization exceeded the set time limit in case of **cube-64** and **k3-2** matrices. Secondly, we noticed the library crashed during processing of **k3-18**, **cube-645** and (partially) **pwr-3d** matrices. A debugging process showed that a segmentation fault occurred in *pdgstrf* function during the numerical factorization phase. We still keep working on this problem together with the PETSc team in order to find a solution.

To complete and perform a fair comparison, an additional flat-MPI test was conducted with a 5-point stencil Poisson matrix with 100000 equations. We think this test can partially allow us to estimate parallel performance of systems like **k3-18**, **cube-645** where SuperLU_DIST crashed. The results of the test are given in figure 1.19.

MPI	MUMPS	PaStiX	SuperLU	MPI	MUMPS	PaStiX	SuperLU
1	7.02E-02	8.72E-02	3.17E+00	11	7.55E-02	8.89E-02	5.82E-01
2	6.73E-02	7.10E-02	1.43E+00	12	7.61E-02	1.06E-01	4.37E-01
3	6.36E-02	7.01E-02	1.07E+00	13	7.84E-02	9.72E-02	5.43E-01
4	6.28E-02	7.11E-02	8.17E-01	14	8.06E-02	1.02E-01	4.22E-01
5	6.50E-02	7.15E-02	7.51E-01	15	8.20E-02	1.19E-01	3.91E-01
6	6.72E-02	7.62E-02	6.15E-01	16	8.07E-02	1.19E-01	4.44E-01
7	6.91E-02	7.69E-02	6.48E-01	17	8.38E-02	1.22E-01	5.19E-01
8	6.89E-02	8.17E-02	5.41E-01	18	8.40E-02	1.26E-01	3.77E-01
9	7.50E-02	8.28E-02	5.02E-01	19	8.58E-02	1.33E-01	5.47E-01
10	7.22E-02	8.52E-02	4.64E-01	20	8.64E-02	1.49E-01	3.39E-01

Table 1.8.: Results of a flat-MPI test of MUMPS, PaStiX and SuperLU_DIST libraries with their default settings and the matrix **cube-5** (9352 equations)

According to the test results, it is clear that MUMPS significantly outperforms both SuperLU_DIST and PaStiX. A literature review showed that Gupta, Koric, and George, in paper [GKG09], came to nearly the same results comparing parallel performance of WSMP, MUMPS and SuperLU_DIST libraries with respect to execution time for their matrix set. However, paper [KBK16] showed an almost opposite outcome. According to Kwack, Bauer, and Koric, SuperLU_DIST spent less time on factorization and solution phases, which almost always determines the total execution time, and, even more interesting, it scaled much better overall. We must note that both research groups used different machines and matrix sets. This fact actually seconds our idea that a choice of a suitable library can depend heavily on data and hardware.

place for
figure
1.19

Taking into consideration the results of our primary flat-MPI tests, the MUMPS library was chosen as a sparse direct solver for our hybrid approach and the following study. Furthermore, an overview of the MUMPS documentation also showed some room for performance tuning that we were going to discuss in detail in sections [bla], [bla] and [bla].

However, it should be mentioned that we cannot exclude that SuperLU_DIST and PaStiX can perform similar, or even better, as the MUMPS library with appropriate parameters tuning or for another matrix set.

MPI	MUMPS	PaStiX	SuperLU	MPI	MUMPS	PaStiX	SuperLU
1	1.36E+00	1.39E+00	time-out	11	7.75E-01	8.15E-01	time-out
2	1.00E+00	9.82E-01	time-out	12	7.81E-01	8.10E-01	time-out
3	8.83E-01	1.06E+00	time-out	13	7.85E-01	8.35E-01	time-out
4	8.17E-01	8.74E-01	time-out	14	7.85E-01	8.18E-01	time-out
5	7.85E-01	8.50E-01	time-out	15	7.88E-01	8.46E-01	time-out
6	8.06E-01	8.52E-01	time-out	16	7.81E-01	8.23E-01	time-out
7	7.71E-01	8.33E-01	time-out	17	6.83E-01	8.49E-01	time-out
8	7.66E-01	8.33E-01	time-out	18	7.96E-01	8.44E-01	time-out
9	7.93E-01	8.35E-01	time-out	19	8.04E-01	8.65E-01	time-out
10	8.07E-01	8.15E-01	time-out	20	6.85E-01	8.87E-01	time-out

Table 1.9.: Results of a flat-MPI test of MUMPS, PaStiX and SuperLU_DIST libraries with their default settings and the matrix **cube-64** (100657 equations)

MPI	MUMPS	PaStiX	SuperLU	MPI	MUMPS	PaStiX	SuperLU
1	1.55E+02	6.44E+01	crashed	11	1.77E+01	3.81E+01	crashed
2	6.28E+01	4.84E+01	crashed	12	1.60E+01	3.75E+01	crashed
3	5.06E+01	5.02E+01	crashed	13	1.42E+01	3.58E+01	crashed
4	4.17E+01	4.50E+01	crashed	14	1.45E+01	3.59E+01	crashed
5	2.52E+01	3.98E+01	crashed	15	1.47E+01	3.57E+01	crashed
6	2.58E+01	4.29E+01	crashed	16	1.41E+01	3.52E+01	crashed
7	2.65E+01	4.30E+01	crashed	17	1.54E+01	3.45E+01	crashed
8	2.59E+01	3.73E+01	crashed	18	1.52E+01	3.31E+01	crashed
9	1.95E+01	4.08E+01	crashed	19	1.52E+01	3.31E+01	crashed
10	1.91E+01	3.81E+01	crashed	20	1.38E+01	3.16E+01	crashed

Table 1.10.: Results of a flat-MPI test of MUMPS, PaStiX and SuperLU_DIST libraries with their default settings and the matrix **k3-18** (1155955 equations)

1.8. Review of MUMPS Library

Originally, MUMPS library was a part of the PARASOL Project. The project was an ESPRIT IV Long Term Research whose main goal was to build and test a portable library for solving large sparse systems of equations on distributed memory systems [Ame+98]. An important aspect of the research was the strong link between the developers of the sparse solvers and the industrial end users, who provided a range of test problems and evaluated the solvers [Ame+02]. Since 2000 MUMPS had continued as an ongoing project and, by the moment of writing, the library have contained almost 5 main releases.




Figure 1.19.: Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and a 5 point-stencil Poisson matrix (1000000 equations)

It was mentioned in section 1.7 that MUMPS is an implementation of the multifrontal method. Hence, MUMPS sequentially performs all three phases: analysis, numerical factorization and solution. The numerical factorization and solution phases were fully described in detail in section 1.5. It is important to examine the analysis phase of MUMPS library because this phase varies from library to library and plays a significant role on parallel performance.

According to the documentation, the MUMPS analysis phases consists of several pre-processing steps:

1. Fill-reducing pivot order
2. Symbolic factorization
3. Scaling
4. Amalgamation
5. Mapping

1) To handle both symmetric and unsymmetric cases, MUMPS performs fill-in reordering based on $A + A^T$ sparsity pattern. The library provides numerous sequential algorithms for reordering such as Approximate Minimum Degree (AMD) [ADD96], Approximate Minimum Fill (AMF), Approximate Minimum Degree with automatic quasi-dense row detection (QAMD) [Ame97], Bottom-up and Top-down Sparse Reordering (PORD) [Sch01], Nested Dissection coupled with AMD (Scotch) [Pel08], Multilevel Nested Dissection coupled with Multiple Minimum Degree (METIS) [KK09]. Additionally, MUMPS can work together with ParMETIS and PT-Scotch which are extensions of METIS and Scotch libraries for parallel execution. MUMPS also provides the user with an automatic choice option where an appropriate reordering algorithm is selected in run-time based on matrix type and size and the number of processors [17b].

2) Sparsity structures of factors L and U are computed during the step, based on permuted matrix A after fill reducing reordering, in order to build the corresponding elimination tree. All computations are performed on a directed graph $G(A)$ associated with the matrix A .

3) At this step, matrix A is scaled in such a way to get absolute values of *one* along the main diagonal and *less than one* for all off-diagonal entries. Scaling algorithms are based on studies described in detail in works [DK99], [DK01] (for the unsymmetric case) and [DP05] (for the symmetric case). This preprocessing step is supposed to improve numerical accuracy and makes all estimations performed during analysis more reliable [17b]. MUMPS also provides an option to switch off scaling or perform it during the factorization phase.

4) During amalgamation, sets of columns with the same off-diagonal sparsity pattern are group together to create bigger nodes, also known as supernodes. The process leads to restructuring of the initial elimination tree to an amalgamated tree of supernodes which is also known as the *assembly tree*. The main purpose of that step is to improve efficiency of dense matrix operations. An example of the amalgamation process is shown in section 1.5.

5) A host process, chosen by MUMPS, creates a pool of tasks where each task can be either a subtree or type 2 or type 3 node (figure 1.16). Then each task is mapped by the host among all available processes in such a way to achieve good memory and compute balance.

Type 1 nodes are grouped in subtrees, according to the Geist-Ng algorithm [GN89], and each subtree is processed only by one single process to avoid the finest granularity,

which can cause high communication overheads.

In case of type 2 nodes, the host process assigns each node to one process, which is called the master, which holds fully summed rows and columns of the node and perform pivoting and factorization of these rows. During the numerical factorization phase, in run-time, the master process first receives symbolic information which describes the structure of the contribution blocks sent by its children. At the next step, the master collects information concerning the load of all other processes and decides which of them (*slaves*) are going to participate. Then the master informs the chosen slaves that a new task has been allocated for them and sends them the frontal matrix distribution. After that, the slaves communicate with the children of the master process and collect the corresponding numerical elements. The slaves are in charge of all the assembly and computation of the partly summed rows.

The root node belongs to the type 3. The host statically assigns the master for the root, as it is in case of type 2 nodes, to hold all the indices describing the structure of the frontal matrix. Before factorization, the structure of the frontal matrix of the root is statically mapped onto a 2D grid of processes using block cyclic distribution. This allows to determine, during the analysis phase, which process an entry of the root is assigned. Hence, the original matrix entries and the part of the contribution blocks can be assembled as soon as they are available. Due to partial pivoting, the master process collects the index information for all delayed variables of its sons, builds the final structure of the root frontal matrix and broadcast the corresponding symbolic information to all slave processes. The slave, in turn, adjust their local data structure. After that numerical factorization can be performed in parallel.

It is important to mention that if the size of the root node is less than a certain computer dependent parameter, defined by MUMPS, the root node will be treated as the type 2.

An illustrative example of process mapping together with a combination of static and dynamic scheduling is given in figure 1.20.

Another outstanding feature of MUMPS is treatment of partial pivoting during the numerical factorization phase. To handle this, MUMPS uses threshold pivoting and delayed pivots approaches which are fully described in section 1.7 where different implementations of direct sparse solvers are compared.

place for
figure
fig:mumps:map
and-
scheduling



Fill reducing reordering is the first and the most important step since it has its direct impact on the assembly tree structure. As we mentioned above, the tree structure defines the task parallelism as well as sizes of frontal matrices and thus performance of the method.

In this section we are going to investigate parallel performance of PT-Scotch and ParMETIS packages as well as MUMPS automatic run-time library selection in case of GRS matrix set. The algorithmic difference between PT-Scotch and ParMETIS was explained in section 1.8.

To perform a test, the default PETSc, MUMPS, PT-Scotch and ParMETIS libraries were downloaded, compiled and configured together. The test was carried out using only flat-MPI mode without explicit process pinning. The results are shown in figure BRA and in appendix BRA-BRA.

As the results show, ParMETIS scales much better in contrast to PT-Scotch. In average, the run-time is reduced by approximately BLA % and BRA % around the saturation point. Unfortunately, it is difficult for us to give a detailed explanation of the results because there is no access to the assembly tree structure from outside of MUMPS, known to us. Additionally, profiling and tracing of both packages, ParMETIS and PT-Scotch, are required which is out of the scope of this study.

1.10. MUMPS: Process Pinning

Due to intensive and complex manipulations with frontal and contribution matrices, we can assume that MUMPS belongs to memory bound applications. In this case memory access can be a bottleneck for the library. A common way to improve performance of memory bound applications running on distributed memory machines is to distribute processes equally among sockets of a node.

However, because MUMPS uses both task and data parallelism as well as a complex hybrid, both static and dynamic, task scheduling, it becomes difficult to decide which pinning strategy is better i.e. *close* or *spread*, described in section 1.2.

Therefore, a couple of tests were conducted with both GRS and SuiteSparse matrix sets in order to investigate which strategy worked better in terms of parallel performance. For this group of tests, we used only MUMPS default settings. The tests were performed on HW1 machine using only flat-MPI mode. Some results are shown in figure 1.21 and in appendix BRA-BRA. During the tests, the total time spent, i.e. analysis, factorization and solution, was measured.

The tests revealed that *spread*-pinning performed better and allowed to reduce run-time by approximately 5.5% in average in contrast to the *close* strategy. As expected, the points with 1 and 20 MPI processes show the same performance because they basically represent the same process distribution. Additionally, almost 13.5% improvement can

place for
figure
1.21

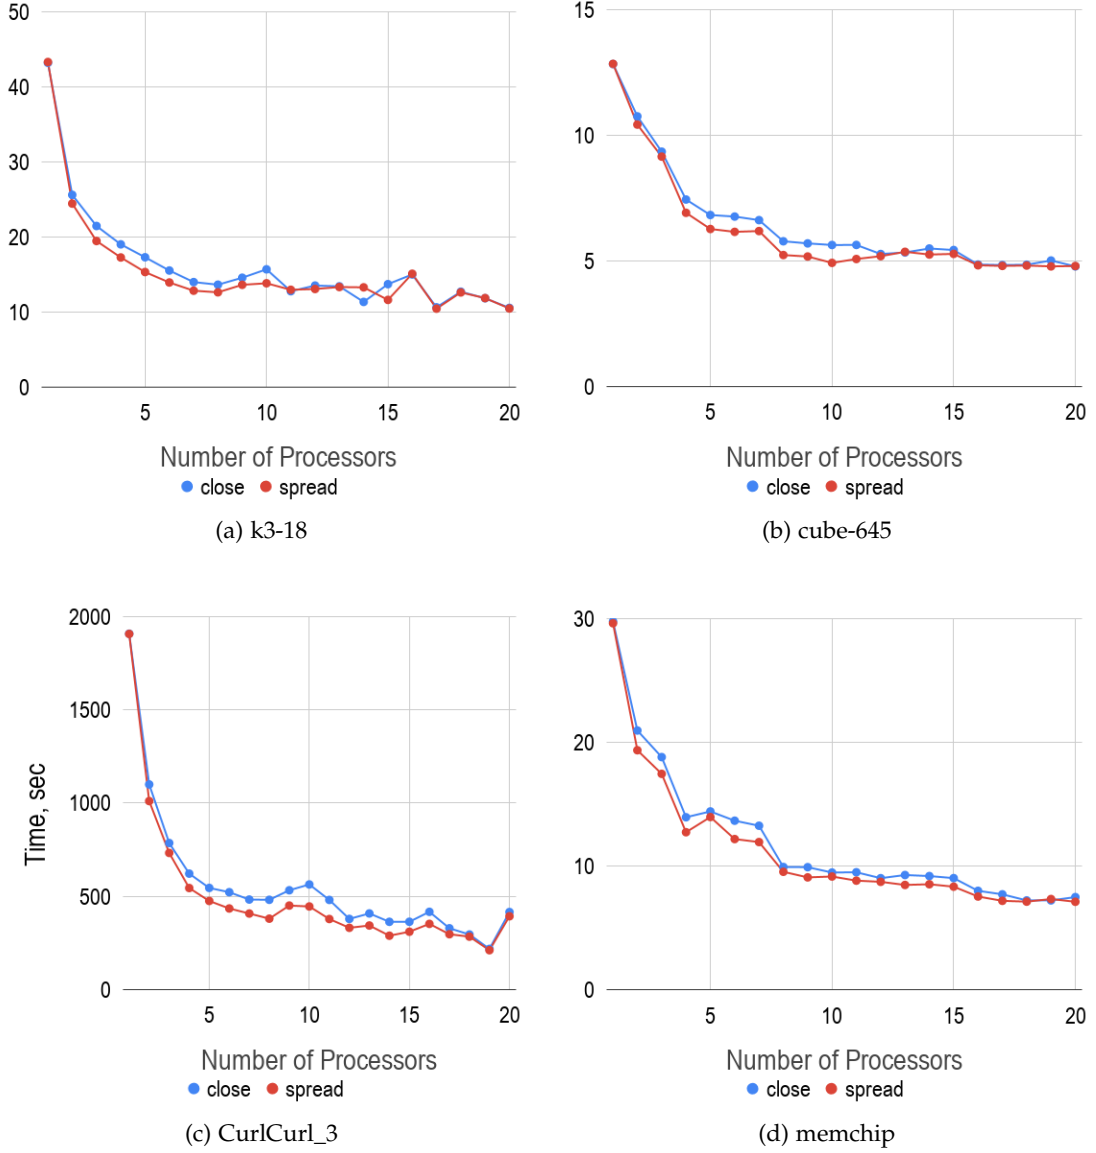


Figure 1.21.: Comparison of *close* and *spread* pinning strategies

be observed around the saturation point in case of *spread* strategy application.

Taken into account test results, *spread*-pinning has been chosen for the rest of the study. Such process distribution can be easily achieved by means of some advanced OpenMPI options, for example *-map-by*, as following:.

```
1 mpiexec -map-by socket -n $num_proc $executable_name $parameters
```

Listing 1.3: An example of *spread*-pinning with using OpenMPI options in case of a flat-MPI run

1.11. Choice of BLAS Library

To perform columns elimination of fully summed block of frontal matrices, MUMPS intensively uses GEMM, TRSM and GETRF subroutines which are parts of BLAS and LAPACK libraries. As an example, figure 1.23 demonstrates factorization of a type 2 node.




Figure 1.22.: MUMPS: static and dynamic scheduling

Both BLAS and LAPACK originate from the Netlib project which is a repository of numerous scientific computing software maintained by AT&T Bell Laboratories, the




Figure 1.23.: MUMPS: An example of a type 2 node factorization

University of Tennessee, Oak Ridge National Laboratory and other scientific communities [06].

The goal of BLAS library is to provide a high efficient implementation of common dense linear algebra kernels by means of high rate of floating point operations per memory access, low cache and Translation Lookaside Buffer (TLB) miss rates.

In its turn, LAPACK is designed in such a way so that as much as possible of computation is performed by calls to BLAS library. This allows to achieve high efficiency for operations such as *LU*, *QR*, *SVD* decompositions, triangular solve, etc. on modern computers. However, the Netlib BLAS implementation is written for an abstract general-purpose central processing unit, in mind, where hardware parameters are based on market statistics. Hence, it is not possible to achieve the maximum possible performance on a specific machine.

Hence, there exist special-purpose, hardware-specific implementations of the library developed by hardware vendors i.e. IBM, Cray, Intel, AMD, etc., as well as open-source tuned implementations such as ATLAS, OpenBLAS, etc. To achieve full compatibility, the developers consider the Netlib implementation of BLAS library as the standard (or reference) and thus overwrite all subroutines with additional tuning and optimization. This approach makes it possible to easily replace different BLAS implementations during object files linking without any modifications of the source code.

Table 1.11 shows commercial and open-source tuned BLAS implementations available on the market today.

Among all libraries listed in table 1.11 there were only four available on HW1 machine, namely: Netlib BLAS, Intel MKL, OpenBLAS and ATLAS. However, installation of ATLAS requires to switch off dynamic frequency scaling, also called CPU throttling, to allow an ATLAS configuration routine to find the best loop transformation parameters for a specific hardware. In order to turn off CPU throttling, one has to reboot the entire machine and make appropriate changes in Basic Input/Output System (BIOS). This fact made ATLAS library not to be suitable for the rest of the research and we excluded it from our primary list of candidates. Moreover, during installation, one has to explicitly provide the number of OpenMP threads that are going to be used once a BLAS subroutine is called. This means there is no way to change the number of threads per MPI process in run-time without re-installation of ATLAS library. Thus, only 3 versions of MUMPS-PETSc (Netlib BLAS, Intel MKL and OpenBLAS) libraries were compiled and tested with using GRS matrix set and 1 thread per MPI process. The test

1. Solver selection and configuration

Name	Description	License
Accelerate	Apple's implementation for macOS and iOS	proprietary license
ACML	BLAS implementation for AMD processors	proprietary license
C++ AMP	Microsoft's AMP language extension for Visual C++	open source
ATLAS	Automatically tuned BLAS implementation	open source
Eigen BLAS	BLAS implemented on top of the MPL-licensed Eigen library	open source
ESSL	optimized BLAS implementation for IBM's machines	proprietary license
GotoBLAS	Kazushige Goto's implementation of BLAS	proprietary license
HP MLIB	BLAS implementation supporting IA-64, PA-RISC, x86 and Opteron architecture	proprietary license
Intel MKL	Intel's implementation of BLAS optimized for Intel Pentium, Core, Xeon and Xeon Phi	proprietary license
Netlib BLAS	The official reference implementation on Netlib	open source
OpenBLAS	Optimized BLAS library based on GotoBLAS	open source
PDLIB/SX	BLAS library targeted to the NEC SX-4 system	proprietary license
SCSL	BLAS implementations for SGI's Irix workstations	proprietary license
Sun Performance Library	Optimized BLAS and LAPACK for SPARC, Core and AMD64 architectures under Solaris 8, 9, and 10 as well as Linux	proprietary license

Table 1.11.: Commercial and open source BLAS implementations [Wik18b]

results are shown in figures bla, bla and bra.

1.12. MPI-OpenMP Tuning of MUMPS Library

As it was mentioned in section 1.8, the development of MUMPS began in 1996 when message-passing programming paradigm dominated in parallel computing. Therefore, the library originally was designed only for distributed-memory machines.

In 2010, Chowdhury and L'Excellent published their first experiments and some issues, in [CL10], of exploiting shared memory parallelism in MUMPS. The authors showed that it was possible to achieve some improvements in multicore systems with multithreading, given a purely MPI application. However, later L'Excellent and Sid-Lakhdar mentioned, in [LS13], that adaptation of existing code for NUMA architecture was still a challenge because of memory allocation, memory affinity, thread pinning and other related issues.

In spite of natural data locality of message-passing applications which is always beneficial, a general motivation of switching to a hybrid mode, mixed MPI/OpenMP, is to reduce communication overheads between processes. According to the profiling results done by Chowdhury and L'Excellent, MUMPS contained four main initial sources of shared-memory parallelization, namely:

1. BLAS Level 1, 2, 3 operations during both factorization and solution phases
2. Assembly operations, where contribution blocks of children nodes of the assembly tree are assembled at the parent level
3. Copying contribution blocks during stacking operations
4. Pivot search operations

Almost all customized BLAS libraries, for example Intel MKL and OpenBLAS, are multi-threaded and can efficiently work in shared-memory environment. Thus, parallelization of region 1 can be achieved by linking a suitable BLAS library whereas regions 2, 3 and 4 are multithreaded by inserting appropriate OpenMP directives above the corresponding loop statements.

A detailed review of works [LS13] and [CL10] reveals that, in general, a pure OpenMP or mixed MPI/OpenMP strategy can reduce run-time of MUMPS. In average, factorization time is reduced by **14.3%** and in some special cases improvement reaches about **50.4%**, according to the data provided in the papers. However, at the same time, the results demonstrate that pure-MPI mode sometimes can significantly outperform any

hybrid mixed strategy.

By and large, the results show two important aspects. Firstly, performance of a specific strategy depends heavily on the resultant assembly tree and thus on the matrix sparsity pattern and applied fill reducing algorithm. Secondly, it is not possible to guess in advance which strategy gives the best parallel performance without detailed information about the tree structure and computational cost per node. L'Excellent and Sid-Lakhdar showed that performance of a particular mode depended on the ratio of large and small fronts. For example, they noticed more threads per MPI process led to better parallel performance when the ratio was high. On the other hand, they observed the absolutely opposite result with relatively small ratios. Unfortunately, L'Excellent and Sid-Lakhdar did not provide any quantitative measure for that in their work [LS13].

It is also interesting to notice that parallelization of region 1 with using a multi-threaded BLAS library brings most of parallel performance for mixed or pure OpenMP strategies, according to the results from [LS13]. But, at the same time, performance of regions 2, 3, 4 multithreaded by OpenMP directives is marginal. In average, it allows to reduce numerical factorization run-time by only **0.66%**.

This outcome is expected because BLAS subroutines, especially level 3, have a high ratio of floating point operations per memory access and thus PEs re-use data stored in caches as much as possible. Meanwhile, regions 2, 3, 4 mainly perform initialization, data movement and execution of *if-statements* which lead to a low compute ratio.

Additionally, it is worth noticing that both works, [CL10] and [LS13], were mainly focused on the numerical factorization phase assuming that both analysis and solution phases do not take lots of time. In spite of credibility of this assumption, it still should be pointed out the solution phase runs faster in case of flat-MPI mode. This fact becomes even more interesting because, in our case, a system with multiple right-hand sides has to be solved in order to generate a preconditioner.

We have to admit that both works, [CL10] and [LS13], are relatively old and the analysis above may be not complete and full. Because MUMPS is a dynamic developing project, we can expect that adaptation of shared-memory parallelization in MUMPS has been significantly advanced since that time. Since the release of MUMPS version 4, the developers have persistently recommended to use only a hybrid mode like *one MPI process per socket and as many threads as the number of cores* [17b].

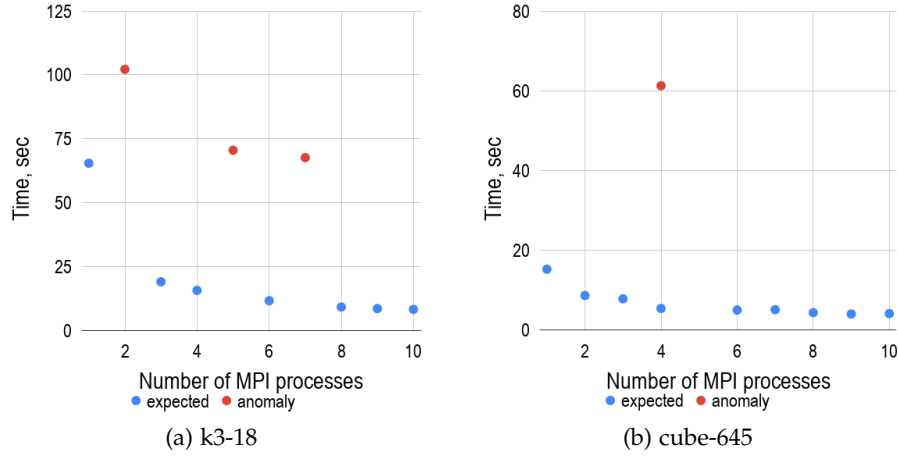


Figure 1.24.: Anomalies of MUMPS-OpenBLAS configuration running with 2 OpenMP threads per MPI process

As an initial test, we decided to compare influence of both Intel MKL and OpenBLAS libraries on parallel performance of MUMPS using GRS matrix set only. In order to pin OpenMP threads in a right way, without any conflict between them, the following OpenMP environment variables were set as:

- OMP_PLACES=cores
- OMP_PROC_BIND=spread

During the test, we found that run-time of MUMPS-OpenBLAS configuration abnormality increased for some test cases. For instance, in case of matrix *cube-645*, the increase reached almost 450% in contrast to the pure sequential execution.

Multiple conflicts between application and system threads were observed using *htop* as an interactive process viewer. Figure 1.25 shows a snapshot taken during factorization of matrix *k3-18* running with 1 MPI process and 20 threads.

place for
figure
1.24

It is difficult to say what exactly caused such behavior. However, Chowdhury and L'Excellent also reported about the same problem with using GotoBLAS (OpenBLAS). They assumed that GotoBLAS created and kept some threads active even after the

place for
figure
1.25




Figure 1.25.: A MUMPS-OpenBLAS thread conflict in case of $k3-18$ matrix factorization (green - application threads, red - system threads)

main threads returned to the calling application which could lead to interference with threads created in other OpenMP regions [CL10]. For that reason, we decided to stick only to Intel MKL library for the rest of the study in this section because there were no conflicts detected during initial runs.

At the beginning, only common mixed MPI/OpenMP modes were tested in order to check influence of shared-memory parallelism on parallel performance of MUMPS as well as to limit the amount of testing. The following strategies were chosen, namely: 20 MPI - 1 thread (flat-MPI), 10 MPI - 2 threads, 4 MPI - 5 threads, 2 MPI - 10 threads, 1 MPI - 20 threads (flat-OpenMP). Additionally, MUMPS was set as a preconditioning algorithm for the GMRES solver with only *one iteration*. This allowed to force MUMPS library to solve systems multiple right-hand sides. According to our assumption, time spent on one GMRES iteration is negligible in contrast time spent on sparse direct matrix factorization. The test results are shown in tables BRA, BRA and BRA as well as in appendix BRA. Numerical values in tables are given in seconds.

1.13. Parallel Performance of MUMPS in case of Small Systems of Equations

1.14. Conclusions

1.15. Outlook

- enhance the GRS matrix set to include more cases. Creation of database or collection
- create automatic testing environment. Since solvers to test solver updates
- Optimization engine for preconditioner parameter search.
- invasive computing (resource-aware programming)

Appendices

A. Choice of a Sparse Direct Solver Library

MPI	MUMPS	PaStiX	SuperLU	MPI	MUMPS	PaStiX	SuperLU
1	4.58E-02	5.60E-02	4.64E+00	11	5.93E-02	8.97E-02	crashed
2	4.31E-02	5.14E-02	1.89E+00	12	6.07E-02	9.20E-02	3.61E-01
3	4.51E-02	5.28E-02	1.22E+00	13	6.26E-02	8.25E-02	crashed
4	4.61E-02	5.64E-02	9.13E-01	14	6.28E-02	9.75E-02	crashed
5	4.92E-02	5.97E-02	7.70E-01	15	6.43E-02	1.03E-01	3.05E-01
6	5.37E-02	6.14E-02	6.04E-01	16	6.55E-02	1.05E-01	2.99E-01
7	5.42E-02	6.51E-02	crashed	17	6.61E-02	9.46E-02	crashed
8	5.41E-02	6.60E-02	4.81E-01	18	6.73E-02	1.24E-01	2.65E-01
9	5.69E-02	6.84E-02	4.35E-01	19	6.84E-02	1.14E-01	crashed
10	5.86E-02	7.22E-02	4.08E-01	20	7.02E-02	1.32E-01	2.60E-01

add information about SuperLU

Table A.1.: Results of a flat-MPI test of MUMPS, PaStiX and SuperLU_DIST libraries with their default settings and the matrix **pwr-3d** (6009 equations)

MPI	MUMPS	PaStiX	SuperLU	MPI	MUMPS	PaStiX	SuperLU
1	1.55E+02	6.44E+01	time-out	11	1.77E+01	3.75E+01	time-out
2	6.28E+01	4.84E+01	time-out	12	1.60E+01	3.58E+01	time-out
3	5.06E+01	5.02E+01	time-out	13	1.42E+01	3.59E+01	time-out
4	4.17E+01	4.50E+01	time-out	14	1.45E+01	3.57E+01	time-out
5	2.52E+01	3.98E+01	time-out	15	1.47E+01	3.52E+01	time-out
6	2.58E+01	4.29E+01	time-out	16	1.41E+01	3.45E+01	time-out
7	2.65E+01	4.30E+01	time-out	17	1.54E+01	3.31E+01	time-out
8	2.59E+01	3.73E+01	time-out	18	1.52E+01	3.31E+01	time-out
9	1.95E+01	4.08E+01	time-out	19	1.52E+01	3.16E+01	time-out
10	1.91E+01	3.81E+01	time-out	20	1.38E+01	3.15E+01	time-out

Table A.2.: Results of a flat-MPI test of MUMPS, PaStiX and SuperLU_DIST libraries with their default settings and the matrix **k3-2** (130101 equations)

MPI	MUMPS	PaStiX	SuperLU	MPI	MUMPS	PaStiX	SuperLU
1	1.52E+01	1.61E+01	crashed	11	8.62E+00	9.09E+00	crashed
2	1.13E+01	1.13E+01	crashed	12	8.53E+00	8.92E+00	crashed
3	1.00E+01	1.03E+01	crashed	13	8.44E+00	9.13E+00	crashed
4	9.29E+00	1.05E+01	crashed	14	8.52E+00	9.00E+00	crashed
5	8.85E+00	9.84E+00	crashed	15	8.54E+00	9.19E+00	crashed
6	8.43E+00	8.99E+00	crashed	16	8.56E+00	9.05E+00	crashed
7	8.64E+00	9.69E+00	crashed	17	8.65E+00	9.12E+00	crashed
8	8.70E+00	9.12E+00	crashed	18	8.62E+00	8.96E+00	crashed
9	8.91E+00	8.94E+00	crashed	19	8.66E+00	9.30E+00	crashed
10	8.76E+00	9.26E+00	crashed	20	8.66E+00	9.16E+00	crashed

Table A.3.: Results of a flat-MPI test of MUMPS, PaStiX and SuperLU_DIST libraries with their default settings and the matrix **cube-645** (1000045 equations)

List of Figures

1.1. An example of process pinning of 5 MPI processes with 2 OpenMP threads per rank in case of HW1	5
1.2. Right-looking parallel LU decomposition	7
1.3. Strong scaling of right-looking <i>LU</i> decomposition using LAPACK and OpenBLAS	8
1.4. A general view on BLAS, LAPACK and ScaLAPACK libraries [17a] . . .	9
1.5. Performance of sparse triangular solve [Jos+98]	11
1.6. GMRES strong scaling speed-up	13
1.7. An example of a sparse matrix and its Cholesky factor [Liu92]	16
1.8. The elimination tree for the matrix example in Figure 1.7 [Liu92]	17
1.9. Information flow of the multifrontal method	19
1.10. An example of matrix postordering from [Liu92]	21
1.11. The stack contents for the postordering [Liu92]	21
1.12. An example of a supernodal elimination tree [Liu92]	23
1.13. Parallel steps of the multifrontal method based on the example in Figures 1.12	24
1.14. Simple parallel models of the multifrontal method	25
1.15. Theoretical speed-up	26
1.16. MUMPS parallelism management in case of 4 PEs [OA07]	28
1.17. Comparison between model 1 and numerical factorization of the matrix <i>cant</i> using MUMPS library	29
1.18. Sparsity structure of the matrix <i>cant</i> before and after fill-in reduction .	29
1.19. Results of a flat-MPI test of MUMPS, PaTiX and SuperLU_DIST libraries with their default settings and a 5 point-stencil Poisson matrix (1000000 equations)	39
1.20. MUMPS: static and dynamic scheduling [LEx12]	42
1.21. Comparison of <i>close</i> and <i>spread</i> pinning strategies	44
1.22. MUMPS: static and dynamic scheduling	45
1.23. MUMPS: An example of a type 2 node factorization	46
1.24. Anomalies of MUMPS-OpenBLAS configuration running with 2 OpenMP threads per MPI process	51

1.25. A MUMPS-OpenBLAS thread conflict in case of $k3-18$ matrix factorization (green - application threads, red - system threads)	52
---	----

List of Tables

1.1. GRS matrix set	3
1.2. SuiteSparse matrix set	3
1.3. Hardware specification	4
1.4. Potential speed-up of linear and quadratic models	26
1.5. Efficiency of linear and quadratic models using 20 PEs	27
1.6. Distribution workload per level in case model 1 and 2	27
1.7. List of packages to solve sparse linear systems using direct methods on distributed memory parallel machines [Li18], [Bal+18]	33
1.8. Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and the matrix cube-5 (9352 equations) . . .	37
1.9. Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and the matrix cube-64 (100657 equations) .	38
1.10. Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and the matrix k3-18 (1155955 equations) . .	38
1.11. Commercial and open source BLAS implementations [Wik18b]	48
A.1. Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and the matrix pwr-3d (6009 equations) . . .	55
A.2. Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and the matrix k3-2 (130101 equations) . . .	55
A.3. Results of a flat-MPI test of MUMPS, PasTiX and SuperLU_DIST libraries with their default settings and the matrix cube-645 (1000045 equations) .	56

Bibliography

- [06] *Netlib Frequently Asked Questions*. 2006. URL: <http://www.netlib.org/misc/faq.html#2.1>.
- [11] *SuperLU Users' Guide*. 2011.
- [17a] *Intel Math Kernel Library Documentation, Overview of ScaLAPACK Routines*. 2017. URL: <https://software.intel.com/en-us/mkl-developer-reference-c-overview-of-scalapack-routines>.
- [17b] *MULTifrontal Massively Parallel Solver, Users' guide*. 2017.
- [AD90] E. Anderson and J. J. Dongarra. "Evaluating Block Algorithm Variants in LAPACK". In: *UT, CS-90-103* (1990).
- [ADD89] M. Arioli, J. Demmel, and I. S. Duff. "Solving sparse linear systems with sparse backward error". In: *SIAM Journal on Matrix Analysis and Applications* (1989).
- [ADD96] P. R. Amestoy, T. A. Davis, and I. S. Duff. "An Approximate Minimum Degree Ordering Algorithm". In: *SIAM Journal on Matrix Analysis and Applications* (1996).
- [Ame+02] P. R. Amestoy, I. S. Duff, J. Koster, and J.-Y. L'Excellent. *MUMPS: A Multifrontal Massively Parallel Solver*. 2002. URL: https://www.ercim.eu/publication/Ercim_News/enw50/amestoy.html.
- [Ame+98] P. R. Amestoy, I. S. Duff, J.-Y. L'Excellent, and P. Plecháč. "PARASOL: An Integrated Programming environment for Parallel Sparse Matrix Solvers". In: *PINEAPL Workshop, A Workshop on the Use of Parallel Numerical Libraries in Industrial End-user Applications*. CERFACS, Toulouse, France, 1998.
- [Ame97] P. Amestoy. "Recent progress in parallel multifrontal solvers for unsymmetric sparse matrices". In: *Proceedings of the 15th World Congress on Scientific Computation, Modelling and Applied Mathematics, IMACS*. Vol. 97. 1997.

- [Bal+18] S. Balay, S. Abhyankar, M. Adams, J. Brown, P. Brune, K. Buschelman, L. Dalcin, A. Dener, V. Eijkhout, W. Gropp, D. Kaushik, M. Knepley, D. May, L. C. McInnes, R. T. Mills, T. Munson, K. Rupp, P. Sanan, B. Smith, S. Zampini, H. Zhang, and H. Zhang. *PETSc Web page*. <http://www.mcs.anl.gov/petsc>. 2018.
- [CCV18] J. Cornelis, S. Cools, and W. Vanroose. “The Communication-hiding Conjugate Gradient Method with Deep Pipelines”. In: *arXiv:1801.04728 [cs.DC]* (2018).
- [CL10] I. Chowdhury and J.-Y. L’Excellent. “Some experiments and issues to exploit multicore parallelism in a distributed-memory parallel sparse direct solver”. PhD thesis. INRIA, 2010.
- [Cor+09] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein. “The recursion-tree method for solving recurrences”. In: *Introduction to Algorithms*. 2009.
- [DGL89] I. S. Duff, R. G. Grimes, and J. G. Lewis. “Sparse Matrix Test Problems”. In: *ACM Trans. Math. Softw.* 15.1 (Mar. 1989), pp. 1–14. issn: 0098-3500. doi: 10.1145/62038.62043.
- [DH11] T. A. Davis and Y. Hu. “The University of Florida Sparse Matrix Collection”. In: *ACM Trans. Math. Softw.* 38.1 (Dec. 2011), 1:1–1:25. issn: 0098-3500. doi: 10.1145/2049662.2049663.
- [DK01] I. S. Duff and J. Koster. “On algorithms for permuting large entries to the diagonal of a sparse matrix”. In: *SIAM Journal on Matrix Analysis and Applications* 22.4 (2001), pp. 973–996.
- [DK99] I. S. Duff and J. Koster. “The design and use of algorithms for permuting large entries to the diagonal of sparse matrices”. In: *SIAM Journal on Matrix Analysis and Applications* 20.4 (1999), pp. 889–901.
- [DP05] I. S. Duff and S. Pralet. “Strategies for scaling and pivoting for sparse symmetric indefinite problems”. In: *SIAM Journal on Matrix Analysis and Applications* 27.2 (2005), pp. 313–340.
- [DR83] I. S. Duff and J. K. Reid. “The Multifrontal Solution of Indefinite Sparse Symmetric Linear”. In: *ACM Transactions on Mathematical Software, (TOMS)* (1983).
- [GG06] K. Goto and R. van de Geijn. *High-Performance Implementation of the Level-3 BLAS*. FLAME Working Note 20. The University of Texas at Austin, 2006.
- [GKG09] A. Gupta, S. Koric, and T. George. “Sparse matrix factorization on massively parallel computers”. In: *IEEE Xplore* (2009).

- [GLU03] A. Guermouche, J.-Y. L'Excellent, and G. Utard. "On the memory usage of a parallel multifrontal solver". In: *Parallel and Distributed Processing Symposium, 2003. Proceedings. International*. IEEE. 2003, 8–pp.
- [GN89] G. Geist and E. Ng. "Task scheduling for parallel sparse Cholesky factorization". In: *International Journal of Parallel Programming* 18.4 (1989), pp. 291–314.
- [Gup00] A. Gupta. *WSMP: Watson Sparse Matrix Package, Part II - direct solution of general systems*. Research Report. IBM T. J. Watson Research Center, 2000.
- [IB70] Irons and Bruce. "A frontal solution program for finite element analysis". In: *International Journal for Numerical Methods in Engineering* (1970).
- [Jos+98] M. V. Joshi, A. Guptam, G. Karypis, and V. Kumar. "A High Performance Two Dimensional Scalable Parallel Algorithm for Solving Sparse Triangular Systems". In: *IEEE Xplore* (1998).
- [KBK16] J. H. Kwack, G. Bauer, and S. Koric. "Performance Test of Parallel Linear Equation Solvers on Blue Waters – Cray XE6/XK7 system". In: May 2016.
- [KK09] G. Karypis and V. Kumar. *MeTis: Unstructured Graph Partitioning and Sparse Matrix Ordering System, Version 4.0*. <http://www.cs.umn.edu/~metis>. University of Minnesota, Minneapolis, MN, 2009.
- [LEx12] J.-Y. L'Excellent. "Multifrontal methods: parallelism, memory usage and numerical aspects". PhD thesis. Ecole normale supérieure de lyon-ENS LYON, 2012.
- [Li18] X. Li. *Direct Solvers for Sparse Matrices*. 2018. URL: <http://crd-legacy.lbl.gov/~xiaoye/SuperLU/SparseDirectSurvey.pdf>.
- [Liu86] J. W. H. Liu. "On the storage requirement in the out-of-core multifrontal method for sparse factorization". In: *ACM Trans. Math. Software* (1986).
- [Liu88] J. W. H. Liu. "Equivalent Sparse Matrix Reordering by Elimination Tree Rotations". In: *SIAM J. Sci. Statist. Comput.* (1988).
- [Liu92] J. W. H. Liu. "The Multifrontal Method for Sparse Matrix Solution: Theory and Practice". In: *SIAM* (1992).
- [LS13] J.-Y. L'Excellent and M. W. Sid-Lakhdar. "Introduction of shared-memory parallelism in a distributed-memory multifrontal solver". PhD thesis. INRIA, 2013.
- [OA07] B. Olivier and G. Abdou. "Task Scheduling for Parallel Multifrontal Methods". In: *Proceedings of the 13th International Euro-Par Conference on Parallel Processing* (2007).

- [Pel08] F. Pellegrini. “Scotch and libScotch 5.1 user’s guide”. In: (2008).
- [PT01] A. Pothen and S. Toledo. “Elimination Structures in Scientific Computing”. In: *CRC Press, LLC* (2001).
- [RBH12] S. Rajamanickam, E. G. Boman, and M. A. Heroux. “ShyLU: A Hybrid-Hybrid Solver for Multicore Platforms”. In: *2012 IEEE 26th International Parallel and Distributed Processing Symposium* (2012).
- [Sch01] J. Schulze. “Towards a tighter coupling of bottom-up and top-down sparse matrix ordering methods”. In: *BIT Numerical Mathematics* 41.4 (2001), pp. 800–841.
- [SS86] Y. Saad and M. H. Schultz. “GMRES: A Generalized Minimal Residual Algorithm For Solving Nonsymmetric Linear Systems”. In: *SIAM Journal on Scientific and Statistical Computing* (1986).
- [Wik18a] Wikipedia contributors. *Amdahl’s law* — *Wikipedia, The Free Encyclopedia*. [Online; accessed 27-November-2018]. 2018.
- [Wik18b] Wikipedia contributors. *Basic Linear Algebra Subprograms* — *Wikipedia, The Free Encyclopedia*. [Online; accessed 14-December-2018]. 2018.
- [Wik18c] Wikipedia contributors. *Runge–Kutta methods* — *Wikipedia, The Free Encyclopedia*. [Online; accessed 10-November-2018]. 2018.
- [Wu12] C. J. Wu. “Partial Left-Looking Structured Multifrontal Factorization and Algorithms for Compressed Sensing”. PhD thesis. University of California, 2012.